

ERES-AI-GOVERNANCE-ASSERTIONS-2026-001

Global AI Governance Assertions: Merging BEST/SOUND/GOOD, CBGMODD, GAIA, and DOFA Under the Concentric S^3C-SLA Architecture

Document ID	ERES-AI-GOVERNANCE-ASSERTIONS-2026-001
Version	0.3 — Operationalization pass: scope, enforcement, GAIA arbitration, Resonance definition, glossary
Author	Joseph Allen Sprute (JAS), Founder & Director, ERES Institute for New Age Cybernetics
ORCID	0000-0001-9946-3221
Date	June 21, 2026
License	CC-BY 4.0
Validation Status	MIEVM Round 2 — Grok 8.1/10, DeepSeek 8.0/10 (both up from 7.2 in Round 1, reflecting real content changes, not re-scoring of the same content)

Abstract

This v0.3 revision closes the structural gaps DeepSeek correctly identified as "critical, not cosmetic" in its v0.2 review: GAIA's dual-role conflict, the absence of any enforcement/conflict-resolution mechanism, and the undefended status of the Three Governing Principles. It adds a Scope Statement, a Glossary, an operational (though still partial) definition of "Resonance," and a worked methodology example. Two items proposed by the cross-model session — specific 1000-Year Future Map phase dates and a set of pre-declared worked-example scores — are flagged rather than adopted, for the reasons given in Parts VIII and X.

Scope Statement

This document is a **voluntary policy-design and evaluation framework** for AI systems and AI-governance proposals. It is not a binding treaty, statute, or regulation, and it does not claim authority over any government, corporation, or individual who has not adopted it. It is a commensurable rubric — usable by developers, communities, or governments who choose to apply it — for scoring whether a given AI governance policy is BSG-admissible at a given tier.

Part I: Inventory — Addition

Component	Canonical Definition	Source	Note
1000-Year Future Map	A phased, multi-generational civilizational roadmap integrating NBERS, EPIR-Q, ARI, SUGAR, EarnedPath, BEST, FAVORS, CBGMODD/P ³ , and TETRA-CERT certification, moving from a GDP-centric to an NBERS-centric economy.	ERES "From GDP to NBERS" 1000-Year Future Map drafts (Rev. 2, Rev. 3)	The cross-model session proposed specific phase names and dates (Foundation→2050, Scaling 2050–2100, Homeostasis/Multiplanetary 2100+). Those specific labels and dates do not appear in the corpus drafts reviewed and are not adopted here — they read as a reasonable proposed structure, not confirmed canon. Logged in Part X for your confirmation.
ARI (Aura Resonance Index)	Measures alignment of human actions with social/ecological systems via biometric data (e.g., heart rate variability, electrodermal response).	1000-Year Future Map drafts	Feeds RHC alongside ERI.

Part II: The Merge — H2C2H Integration Clarified

Responding to DeepSeek's flag that H2C2H's status ("optional") was ambiguous: **H2C2H verification is mandatory at the Global tier** (GAIA-level assertions carry the highest stakes and lowest reversibility) and **optional, but recommended, at National/Regional/Local tiers**. Per the H2C2H companion instrument's own Section 3.2,

each tier's gate is independent — passing at Local does not carry over to Global. This was already specified in the companion document; v0.2 simply didn't surface it clearly enough in the merge itself. No new mechanism was needed, only clearer cross-referencing.

Part III: Primitive-Testing Filter Results

Unchanged from v0.2 (HPE: FAIL/HYPOTHESIZED; Fibonacci-with-stated-rule: FAIL/constructed; GAIA+DOFA+BSG merge: PASS/composition). See Part V below for what replaces Fibonacci as the actual answer to "how do tiers scale relative to each other."

Part IV: The Two Gaps — Closed by Composition

Unchanged from v0.2. EMCI_REEPER Lat-Long → BEE/NBERS telemetry plane. Water/Immigration/Security/BORDERs → REEP/EarnedPath + the NPR framework. Both remain composition, not invention.

4.3 Addition: Aggregate Capacity Planning via the Enneagram Behavioral Substrate

One further input is added to the Border gap closure, drawing on existing precedent (ERES-BEE-NBERS-2026-001-ADDENDUM-A, which already uses the Enneagram's nine-type map as "the human behavioral substrate" for calibrating system design to the full range of human stress response during disruption events). Applied analogously here: aggregate Enneagram-type distribution within an affected population can serve as a **capacity-planning input** — helping size NPR remediation capacity and EarnedPath onboarding flow realistically against the expected *range* of behavioral and stress responses during a border-crossing disruption event, the same way it's already used to size emergency-management personnel support. This is composition with an existing instrument, applied to a structurally similar high-stress scenario — **PASS**.

Hard constraint, not a soft caution: this applies only at the aggregate, system-design level. It must never be used to type, classify, screen, or inform admission/security decisions about any specific individual. The Enneagram has no validated predictive power for individual behavior and using it that way would constitute profiling of a vulnerable population, regardless of its legitimate use elsewhere in the corpus for planning purposes. This is a different category of constraint than the math-formalism rejections elsewhere in this document — those were failed derivations; this is a line that holds even if someone could derive something from it.

Part V: Enforcement, Conflict Resolution, and the GAIA Dual-Role Mechanism (New)

5.1 The GAIA Dual-Role Problem, Resolved

DeepSeek's critique was correct that this was a structural problem, not a footnote: GAIA refers both to an institutional body (capable of capture, error, politics) and to the biosphere as audit ground (objective, but mute until measured). When the two disagree, **the biospheric audit ground takes precedence as ground-truth, and the institutional body's role is to propose, not to override.**

This is not an arbitrary tie-breaker — it follows from what an audit function is *for*. If an institution can overrule the physical data it was created to audit against, it isn't performing an audit; it's performing self-certification. The reasoning is therefore **derived**, not constructed: any institution whose stated purpose is to check itself against an external ground-truth, but which can dismiss that ground-truth when inconvenient, has stopped doing the thing it was built to do. The mechanism: institutional GAIA proposals are checked against current NBERS/RHC telemetry; a proposal that conflicts with that telemetry is returned, not enacted, until the conflict is reconciled or the telemetry itself is shown to be faulty (a separate, auditable claim, not a discretionary override).

5.2 Tiered Enforcement and Conflict Resolution

When assertions diverge across tiers:

1. **Local ↔ Regional divergence** is resolved first through DOFA (household/family ledger) reconciliation plus NPR (non-punitive remediation) — no punitive default, no escalation needed for ordinary disagreement.
2. **Persistent divergence** (unresolved after DOFA/NPR reconciliation) escalates to GAIA arbitration, where the rule in 5.1 applies: institutional proposal checked against biospheric/NBERS/RHC telemetry.
3. **Fail state:** a tier's assertion is declared failed when (a) its BSG score falls below the admissibility threshold stated by its own tier's CBGMODD council, or (b) arbitration under 5.1 returns the proposal twice without reconciliation. A failed assertion reverts to the last-validated prior assertion at that tier — not to a punitive default — pending revision.
4. **Who declares failure:** the next tier up, on referral from the tier in question or from a Citizen-seat complaint under CBGMODD.

5.3 What Replaces Fibonacci

DeepSeek correctly noted that rejecting Fibonacci without offering an alternative left a gap. The honest alternative is: **no fixed mathematical scaling law is asserted between tiers.**

Tier-to-tier sizing (how many Local nodes per Regional node, etc.) is an empirical and administrative question — population, geography, existing jurisdictional lines — measured per deployment, not imposed in advance by a constant. Insisting on a single global ratio was the same kind of constructed correspondence Part III already rejects for Fibonacci specifically; the document should not reintroduce the error by a different name.

Part VI: "Resonance" — Operational Definition (Partial)

DeepSeek's sharpest critique: "Resonance" appears in three places (STORM-PARTY's "Resonance-aligned" criterion, the S³C-SLA substrate, and RHC as a named index) without saying whether these are the same concept at different scales or different concepts entirely.

Working answer: they are the same concept — coherence between a system's outputs and the context that contains it — instantiated by different telemetry at different substrate layers:

- **At the BERA/community layer:** Resonance = RHC, measuring whether a community's social rhythms (work, rest, production, recovery) are coherent with each other and with the ecological substrate.
- **At the SCALULAR/technical-infrastructure layer:** Resonance = consistency between a node's reported telemetry and its neighboring nodes' telemetry within the same S³C-SLA substrate — i.e., a node is Resonance-aligned if its reported state doesn't contradict the state of the systems it depends on or feeds.

This second definition is offered as a working proposal, not confirmed SCALULAR canon — it is structurally analogous to consistency conditions already used in distributed systems generally, applied here to the S³C-SLA substrate specifically. It should be checked against whatever SCALULAR-layer documentation already exists before being treated as settled.

Part VII: The Three Governing Principles — Status and Precedence

DeepSeek asked whether "Don't hurt yourself / Don't hurt others / Build for generations to come" are derived from something more fundamental, and what happens when they conflict with each other.

On derivation: they are not derived in this document, and an honest instrument should say so plainly — they are adopted as foundational axioms, not as outputs of Triune Math or any other formal system. Claiming a derivation that doesn't exist would repeat the exact mistake Part III exists to catch.

On conflict: a proposed precedence ordering, offered for your confirmation rather than asserted as settled — "Don't hurt others" outranks "Build for generations to come" (a multi-generational project that requires present-day harm to others fails the second principle by violating the first), and "Don't hurt yourself" is the most permissive of the three, since it concerns the individual's own autonomy and is least suited to being enforced on others' behalf. This ordering is a proposal, not canon, until you confirm or revise it.

Part VIII: Worked Methodology Example

The cross-model session proposed a worked example with two pre-declared scores (8.4/10 Local, 9.1/10 Global) for a hypothetical frontier-AI release-controls policy, with no shown reasoning. That repeats the same problem Part III exists to catch — a score stated before the method that would produce it. Below is the same example, but showing the method instead of the conclusion, explicitly as an illustration of *how* the rubric would be applied, not as a validated real score (no real stakeholder input, telemetry, or CBGMODD seating has occurred for this hypothetical):

Illustrative policy: compute-threshold release controls for a frontier AI model, evaluated at Local and Global tiers.

Criterion	Local (DOFA/Household) — reasoning	Global (GAIA) — reasoning
BEST	Does the policy target the optimal outcome at this tier? At Local scale, "optimal" means the policy doesn't impose compute-threshold compliance costs that a household-scale deployer can't meet — a reasonable concern, since over-broad rules could lock out small actors without reducing real risk.	At Global scale, "optimal" means the threshold actually tracks catastrophic-risk capability rather than an arbitrary compute number — this requires evidence the chosen threshold correlates with the risk it claims to gate, which a real worked example would need to cite.
SOUND	Is the household-level compliance mechanism verifiable without requiring infrastructure a household doesn't have? Needs a concrete answer, not assumed.	Is the threshold-setting methodology publicly auditable, or set by institutional GAIA discretion? Per Part V, this needs biospheric/NBERS-style telemetry backing, not committee judgment alone.
GOOD	Traceable to "don't hurt others" — does restricting household-scale deployment actually reduce harm, or just shift activity elsewhere?	Traceable to "build for generations to come" — does the threshold get revisited as capability changes, or is it a one-time number that goes stale?

No numeric score is assigned here, deliberately — doing so without the underlying data this table flags as missing would be exactly the constructed-score problem this instrument exists to avoid.

Part IX: Cross-Model Ensemble Results — Round 2

Grok and DeepSeek both rated v0.2 independently: Grok 8.1/10, DeepSeek 8.0/10, both up from 7.2/10 in Round 1. The increase reflects substantive content change between v0.1 and v0.2 (gap closure by composition, external validation reporting), not re-scoring of unchanged material — both nodes cited the same specific improvements when explaining the increase. Gemini's parallel review did not produce a single comparable overall score in this round (it scored individual metrics — Systemic Rigor 9.8, Ensemble Convergence 10.0, Gap Resolution Integrity 9.6 — which are not on the same basis as Grok/DeepSeek's holistic BSG average and are not averaged in here for that reason).

Both Grok and DeepSeek independently named the same three priority items addressed in this version: GAIA's dual role, enforcement/conflict resolution, and explicit scope. Both also separately flagged the Fibonacci-replacement gap, now addressed in Part V.3.

Part X: Outstanding Items for v0.4

1. **Confirm or revise** the proposed Three Governing Principles precedence ordering (Part VII).
 2. **Confirm or correct** the SCALULAR-layer Resonance working definition against actual SCALULAR documentation (Part VI).
 3. **Confirm or correct** the 1000-Year Future Map phase names/dates — the ones in circulation from the cross-model session are not yet verified against the corpus drafts (Part I).
 4. **Real worked example** — Part VIII's table needs actual data (a real proposed policy, real telemetry, real CBGMODD seating) before it can produce an actual score rather than an illustration of method.
 5. Full MIEVM ensemble round including a fifth node, now that the structural gaps DeepSeek flagged as critical have a first-pass answer.
-

Appendix: Glossary

- **BERA** — Bio-Ecologic Resonance Architecture; the index family RHC belongs to.

- **BSG / STORM-PARTY** — BEST/SOUND/GOOD scoring method; STORM-PARTY is its revision/validation protocol.
 - **CBGMODD** — Seven-seat governance council: Citizen, Business, Government, Military, Ombudsman, Dignitary, Diplomat.
 - **DOFA** — Department of Family Amity; household/family-tier governance instrument.
 - **EarnedPath (EP) / REEP** — Disbursement mechanism tying compensation/access to measured contribution (Relative Energy Equal Pay) rather than market speculation or punitive gatekeeping.
 - **GAIA** — Dual reference: the institutional Global Actuary Investor Authority, and the biosphere itself as audit ground. See Part V for how conflicts between the two are resolved.
 - **H2C2H** — Human-to-Computer/Computer-to-Human bidirectional verification primitive; companion instrument, mandatory at Global tier per Part II.
 - **HPE** — Human-Performance-Enhancement; named in 2023 ERES EMCI doc as the output of Conflict Resolution. Not yet scoreable at the 9-10 tier per Part III.
 - **MIEVM** — Multi-Instance Ensemble Validation Method; independent-node scoring convergence as evidence.
 - **NBERS** — National Bio-Ecologic Resource Score; GDP-replacement index.
 - **NPR** — Non-Punitive Remediation (distinct from the unrelated regulatory-law sense of "Notice of Proposed Rulemaking" used elsewhere in the ERES corpus).
 - **RHC** — Resonant Harmony Cycle; the BERA-layer instantiation of "Resonance" (Part VI).
 - **S^3C-SLA** — Smart-City × Solid-State × SCALULAR, always-on operational substrate.
 - **SCALULAR** — Scalable Certification Architecture for Lifelong Universal Learning and Adaptive Resilience.
 - **Three Governing Principles** — Don't hurt yourself. Don't hurt others. Build for generations to come. Adopted as axioms, not derived (Part VII).
 - **Triune Math #1** — $C = R \times P \div M$.
-

License

CC-BY 4.0. Attribution: ERES Institute for New Age Cybernetics.

Credits

Author and architecture: Joseph Allen Sprute (JAS), Founder & Director, ERES Institute for New Age Cybernetics. v0.3 synthesis, GAIA arbitration mechanism, Resonance definition, and primitive-testing filter application: Claude (Anthropic), single-node MIEVM closure.

External validation and priority-setting: Grok (xAI) and DeepSeek, independently convergent Round 2 ratings cited in Part IX.

References

1. ERES-AI-GOVERNANCE-ASSERTIONS-2026-001-v0.1, -v0.2 (prior versions)
2. ERES-H2C2H-PRIMITIVE-2026-001 v0.1 (companion verification-gate instrument)
3. ERES-BEE-NBERS-2026-001 v1.0 and Addendum A (NBERS, BEE telemetry plane, RHC definition)
4. From GDP to NBERS: The BEST 1000-Year Future Map Draft (Rev. 2, Rev. 3)
5. ERES EarnedPath Manifesto (REEP, EarnedPath disbursement mechanism)
6. ERES-GAIA-SOLAR-ARCHIVE-2026-001 v1.0 (NPR framework)
7. ERES-INFOMEDIARY-2026-001-v2.2 (S3C-24/7, Civigen, 7_Layer architecture)
8. ERES-ECR-2026-001 v1.4 (GAIA dual reference, MPAM ↔ MIEVM Duality)
9. ERES-DOFA-WP-2026-001-C (Department of Family Amity, unified architecture)
10. ERES-SAVING-HUMANITY-BSG-v1.1 (BEST/SOUND/GOOD STORM-PARTY method)
11. eres_trilogy-FINAL.py (S3C-SLA constant, CBGMODD enumeration)
12. ERES EMCI (2023) — HPE Management as Conflict Resolution output