

ERES INSTITUTE FOR NEW AGE CYBERNETICS

Whitepaper

Erasure Completeness Ratio (ECR): The Falsifiability Metric of MPAM↔MIEVM Duality

Document ID: ERES-ECR-2026-001**Author:** Joseph A. Sprute (JAS), ERES Maestro**Affiliation:** ERES Institute for New Age Cybernetics, Bella Vista, Arkansas**ORCID:** 0000-0001-9946-3221**Date:** 21 May 2026**Version:** v1.0 (Initial deposit)**License:** CARE Commons Attribution License v2.1 (CCAL v2.1)**Series:** ERES Trilogy support corpus — Volume Zero (Infomediary) instrumentation layer

Abstract

The Erasure Completeness Ratio (ECR) is the empirical falsifiability metric of ERES governance. It measures, for any given operational window, the fraction of attestations, permissions, claims, or system states that should have been revoked or retired and that actually were revoked or retired — across every downstream surface that ever held them. ECR emerges directly from the MPAM↔MIEVM Duality: every affirmative attestation has a mirror erasure obligation, and the integrity of the affirmative side is only as strong as the closure of the erasure side. This whitepaper defines ECR formally, anchors it within the canonical Triune Math root $C = R \times P \div M$, specifies its instrumentation through EAAP and EAAP-CRL under ML-DSA/Falcon post-quantum signatures, and demonstrates why ECR (not narrative coherence, not symbolic elegance) is the metric that converts ERES from an abstract narrative system into a governable, empirically testable control law over a one-thousand-year horizon.

1. Background and Provocation

ERES has been challenged — appropriately and repeatedly — on whether its symbolic apparatus is empirically validatable, or whether the equations, backronyms, and lock-boxes constitute an abstract narrative system indistinguishable from rhetoric. The most recent and most precise form of this challenge, raised in May 2026 by an Execution Governance Strategist working on runtime oversight for autonomous AI systems, is reproduced in paraphrase:

The issue is less the symbolic equation itself and more whether the underlying governance logic can be empirically validated. Without definitional clarity around the variables and control structure, these models become difficult to distinguish from abstract narrative systems.

This challenge is legitimate. It is also, within ERES, already answered — but the answer requires naming the specific instrument that closes the loop, rather than restating the equation. That instrument is ECR. The remainder of this document specifies it.

2. The Canonical Root

ECR cannot be defined in isolation. It is the falsifiability metric of a specific control law, and that control law must be stated first, in its canonical form, with no variable migration.

2.1 Triune Math Key #1

$$C = R \times P \div M$$

$$\text{Cybernetics} = \text{Resource} \times \text{Purpose} \div \text{Method}$$

The variables are fixed. Drift in any of them — substituting Consequence for Cybernetics, Permissions for Purpose, Mitigation for Method — produces a different equation wearing the same letters, and is itself a failure mode (repeatable drift) that ECR is partly designed to detect across instruments.

2.2 The Three Governing Principles

Every ECR computation is subordinated to:

1. Don't hurt yourself.
2. Don't hurt others.
3. Build for generations to come.

An ECR of 1.0 obtained at the cost of violating any of these three is, by ERES doctrine, a null result.

2.3 Administration

$C = R \times P \div M$ is administered by the User-GROUP Community of Interest (COI) operating through GAIA, where GAIA names the biospheric substrate — the planet itself as audit ground — and HowWay is its operational handle. ECR is the number that tells the COI whether HowWay is actually doing what HowWay claims to do.

3. The MPAM↔MIEVM Duality

Every ERES governance act has two faces:

MPAM (affirmative side)

Multi-Party Attestation Matrix — what is positively recorded: signed credentials, granted permissions, deposited claims, ratified states, on-chain commitments. MPAM is the corpus of things ERES has said yes to.

MIEVM (validation side)

Multi-Instrument Ensemble Validation Mean — the cross-instrument check, currently realized through a short-list ensemble (Claude, Grok, DeepSeek, ChatGPT) culled by bandwidth constraint from a larger original set. MIEVM tests whether the MPAM record holds up under independent re-derivation.

The duality claim

MPAM is only as trustworthy as MIEVM can re-derive it, and MIEVM is only as trustworthy as the erasure side of MPAM is complete. An attestation that cannot be cleanly revoked is an attestation that was never really under control. The duality is therefore not a redundancy — it is a closure requirement. ECR is the metric of that closure.

4. Erasure Completeness Ratio — Formal Definition

4.1 The ratio

For an operational window W and a population of governance objects O (attestations, permissions, claims, states, credentials):

$$\text{ECR}(W, O) = \frac{|\{o \in O : o.\text{required_erasure}(W) \wedge o.\text{actually_erased}(W)\}|}{|\{o \in O : o.\text{required_erasure}(W)\}|}$$

ECR is bounded in $[0, 1]$. By construction:

- **ECR = 1.0** — every object that should have been erased, was erased, everywhere it propagated.
- **ECR < 1.0** — residual entries persist: stale signatures still verifying, revoked permissions still being honored, retracted claims still propagating downstream. The complement $(1 - \text{ECR})$ is the leakage fraction.
- **ECR = VEILED** — reserved state for windows in which the required-erasure population is itself unresolvable (the denominator cannot be determined). VEILED replaces the earlier $\perp$ (bottom) notation; it is not zero, not undefined, and not failure — it is the

explicit acknowledgement that the question is not currently answerable, which itself is information.

4.2 The seven RATE dimensions

ECR is not a single scalar in production deployments. It is computed across the seven RATE dimensions yielded by the canonical relation $(CBG\text{MODD} \times \text{FAVORS}) + \text{BERA} = \text{RATE}$, indexed R_1 through R_7 . A deployment reports a vector $\langle \text{ECR}_1, \text{ECR}_2, \dots, \text{ECR}_7 \rangle$, and the convergent mean across the seven is what is compared to the MIEVM GO threshold. A scalar reported without the seven-vector behind it is presumptively VEILED.

4.3 Why this is empirically testable

Each term on the right-hand side is operationally defined and externally measurable:

- `required_erasure(W)` is determined by the COI charter and the EAAP-CRL revocation list — both public, both signed, both auditable.
- `actually_erased(W)` is determined by polling every downstream surface (registries, mirrors, cached attestations, replica nodes) and confirming that signatures over revoked objects no longer verify.
- The denominator is finite and enumerable within W ; the numerator is a subset of the denominator by construction.

This is the property that distinguishes ECR from a narrative claim. A narrative claim that “the system retires obsolete permissions” cannot be falsified. ECR can be falsified, in principle, by any third party with access to the EAAP-CRL and the downstream surfaces; if their independent computation diverges from the ERES-reported ECR, the discrepancy is itself a measurable error term.

5. Instrumentation

ECR is not a calculation performed in isolation; it sits on a stack of named ERES instruments, each of which is itself externally inspectable. The table below maps each sub-component of the ECR computation to the instrument that produces it and to the mode by which a third party can falsify the reported value.

ECR sub-component	ERES instrument	Falsifiability mode
<code>required_erasure(W)</code> enumeration	EAAP-CRL revocation list (post-quantum signed)	Third-party verifies CRL signature, recomputes population

ECR sub-component	ERES instrument	Falsifiability mode
actually_erased(W) verification	Downstream surface polling under ML-DSA / Falcon	Independent re-verification of signature non-validity
Subject anchoring	FAVORS six-factor biometric (Fingerprint, Aura, Voice, ODOR, Retina, Signature)	Per-subject erasure traceable to a unique embodied claimant
Cross-instrument check	MIEVM ensemble (Claude, Grok, DeepSeek, ChatGPT) at convergent mean ≥ 8.5	Divergent ECR across instruments → repeatable-drift flag
Tuning signal	BERA Resonance Metrics: ARI, ERI, RHC (Resonant Harmony Cycle), RCI	ECR feeds Meritcoin under Proof-of-Resonance (“tuning, not mining”)
Unresolvable state	VEILED (supersedes $\perp$)	Explicit non-answer; itself audit information

Each row is independently auditable. The full ECR cannot be fabricated without simultaneously fabricating every instrument in the right-hand column — which is precisely the property that converts the system from narrative to control law.

6. ECR as the Falsifiability Hook

The provocation in Section 1 reduces, formally, to: “provide an operationally defined quantity Q such that an external party can compute Q' from the same inputs and compare Q' to ERES-reported Q .” ECR is that quantity.

6.1 The independent re-derivation test

For any window W , an external auditor with access to:

- the published EAAP-CRL for W ,
- the set of downstream surfaces declared in the COI charter,
- the post-quantum verification keys (ML-DSA / Falcon public keys),

can compute their own $ECR'(W, O)$ without any privileged cooperation from ERES. The discrepancy $\Delta = |ECR - ECR'|$ is itself a measurable error term, and a non-zero Δ is a falsification event — either an ERES reporting error, an auditor methodology error, or a substrate compromise. All three are actionable.

6.2 The repeatable-drift detector

A particularly important class of failure that ECR detects: repeatable drift. If across multiple windows $W_1, W_2, \dots, W_n$ the seven-vector $\langle \text{ECR}_1 \dots \text{ECR}_7 \rangle$ trends monotonically downward in any dimension without an explanatory event, the system is leaking erasure in a structured way — quietly accumulating residual permissions or stale attestations. This is the AI-determinism failure mode (deterministic but drifting) reframed in governance terms, and it is distinct from ordinary random error. ECR makes it visible.

6.3 The narrative-system exclusion

A pure narrative system cannot produce an ECR because it has no enumerable required_erasure population. The denominator does not exist. Any system that publishes an ECR (with the seven-vector and the instrument provenance) has, by that publication alone, exited the narrative-system category. ECR is therefore the bright line.

7. ECR Under Post-Quantum Conditions

ECR is meaningful only as long as the signatures it depends on remain meaningful. EAAP migrated to ML-DSA and Falcon in early 2026 specifically to ensure that erasure semantics remain falsifiable on a one-thousand-year horizon (per the Equity Covenant, ERES-COVENANT-2026-001).

Under classical signature schemes, a sufficiently advanced adversary could forge revocation entries or, conversely, forge proofs that revoked objects remain valid. Either attack collapses ECR — the metric returns plausible-looking numbers while the underlying erasure surface is being silently rewritten. ML-DSA (for general attestation) and Falcon (for size-constrained CRL entries) push the falsifiability horizon out beyond foreseeable cryptographic advances, including those leveraging quantum computation.

This is not theoretical hardening. It is the load-bearing requirement for ECR to remain a 1000-year metric rather than a 25-year metric.

8. ECR and Proof-of-Resonance

Meritcoin operates under Proof-of-Resonance, not Proof-of-Work. The distinction matters for ECR specifically:

Proof-of-Work has no native erasure semantics

Once a block is buried under sufficient subsequent work, it is buried. There is no operational notion of “retiring” that block — the chain treats it as permanent. An ECR-equivalent computed against a PoW substrate would always be near zero in the denominator-defined sense (almost nothing is ever required to be erased) and therefore vacuous.

Proof-of-Resonance is continuously revisable

Resonance state is a tuning, not a settlement. The four BERA Resonance Metrics — ARI (Aggregate Resonance Index), ERI (Ethical Resonance Index), RHC (Resonant Harmony Cycle), RCI (Relational Coherence Index) — are re-evaluated at every Meritcoin update window. A prior tuning that no longer holds is not buried; it is retired, and the retirement is exactly what ECR measures. “It’s not mining — it’s tuning” is therefore not a slogan but a structural claim about the erasure surface.

ECR is the metric that audits whether the tuning system is honoring its own retirement obligations. Without ECR, Proof-of-Resonance would have the same auditability gap as Proof-of-Work; with ECR, it has the closure property that PoW lacks by construction.

9. The One-Thousand-Year Horizon

The ERES Equity Covenant (ERES-COVENANT-2026-001) is a one-thousand-year governance instrument. ECR is the metric that makes that horizon empirically meaningful rather than aspirational. Three properties together produce this:

4. Each ECR computation is bounded in a specific window W ; the horizon is composed of contiguous windows, each independently verifiable.
5. The instrumentation (EAAP-CRL under ML-DSA/Falcon, FAVORS biometric anchoring, MIEVM cross-instrument check) is designed to remain falsifiable across substrate generations — cryptographic, computational, and institutional.
6. VEILED windows are explicitly preserved as VEILED, not retroactively coerced into numerical values. A 1000-year record will contain VEILED entries; their preservation is itself part of the audit substrate.

The horizon, in other words, is not maintained by claiming completeness; it is maintained by maintaining the apparatus that makes incompleteness visible.

10. MIEVM Validation Protocol for This Whitepaper

In accordance with the ERES Force-Multiplier Method v1.5, this document is subject to MIEVM cross-instrument validation prior to canonical deposit. The GO threshold is a convergent mean ≥ 8.5 across the four instruments (Claude, Grok, DeepSeek, ChatGPT), with attention to:

- Canonical preservation: variables in $C = R \times P \div M$ must not migrate.
- Subject-First Nomenclature: ECR is named for what it measures (erasure completeness), not for what it serves.

- Falsifiability: every numerical claim must be reproducible by an external party from declared inputs.
- Drift detection: the seven-vector $\langle \text{ECR}_1 \dots \text{ECR}_7 \rangle$ across windows must be preserved, not collapsed to scalar.

Convergence below threshold triggers revision. VEILED is permitted as an instrument response and does not by itself block deposit, provided the VEILED reasoning is itself reported.

11. Conclusion — The Loop Closes

The Erasure Completeness Ratio is the specific empirical instrument that converts ERES governance from a symbolic apparatus into a control law. It does this by making one question operationally answerable across any deployment window and any governance object class:

“When this system says something is gone, is it actually gone, everywhere?”

ECR returns a number in $[0, 1]$, or returns VEILED with reasons. Either return is information. Either return is auditable. Either return distinguishes ERES from an abstract narrative system, by the simple test that no narrative system can produce either return at all.

The root holds: $C = R \times P \div M$, Cybernetics = Resource $\times$ Purpose $\div$ Method, administered by User-GROUP COI through GAIA = HowWay. ECR is the metric by which HowWay demonstrates, window by window, that it is actually doing what it says it does. With ECR specified and instrumented, the loop closes.

Appendix A — Canonical Cross-References

- ERES Triune Math Keys (1) $C = R \times P \div M$; (2) $M \times E + C = R$; (3) $\text{REAL} = (E \cdot M \cdot R) / (T \cdot S)$.
- ERES Equity Covenant — ERES-COVENANT-2026-001 (1000-year governance instrument).
- ERES Attestation and Authorization Protocol — EAAP v1.5, with EAAP-CRL under ML-DSA / Falcon.
- ERES Reliability Whitepaper — ERES-RELIABILITY-2026-001 (five AI failure vectors mapped to ERES instruments).
- ERES Force-Multiplier Method v1.5 — Subject-First Nomenclature; MIEVM GO threshold convergent mean ≥ 8.5 .
- ERES Infomediary (Volume Zero of the ERES Trilogy Architecture), consolidated v2.0.
- UBIMIA Manual v1.1 (25 April 2026) — partial instruction set for the PlayNAC engine.

- FAVORS six-factor biometric: Fingerprint · Aura · Voice · ODOR · Retina · Signature.
- BERA Resonance Metrics: ARI · ERI · RHC (Resonant Harmony Cycle) · RCI.
- CBGMODD seventh seat = Diplomat (not a repeated Dignitary).
- VEILED — supersedes $\perp$ (bottom) as the semantically unresolvable RATE state.

Appendix B — Notation

- **ECR(W, O)** — Erasure Completeness Ratio over window W and object population O .
- $\langle \mathbf{ECR}_1 \dots \mathbf{ECR}_7 \rangle$ — ECR seven-vector across the RATE dimensions $R_1 \dots R_7$.
- **VEILED** — explicit unresolvable state; replaces $\perp$.
- $\Delta = |\mathbf{ECR} - \mathbf{ECR}'|$ — audit divergence between ERES-reported and externally-recomputed ECR.

License and Deposit

This whitepaper is released under the CARE Commons Attribution License v2.1 (CCAL v2.1). Attribution: Joseph A. Sprute, ERES Institute for New Age Cybernetics, Bella Vista, Arkansas, 21 May 2026. ORCID: 0000-0001-9946-3221. Intended deposit: ResearchGate, with concurrent reference at the ERES Library and the ERES Institute GitHub repositories.

Canonical document identifier: ERES-ECR-2026-001.

— **Joseph A. Sprute, ERES Maestro**

ERES Institute for New Age Cybernetics