

ERES INSTITUTE FOR NEW AGE CYBERNETICS

Whitepaper

Erasure Completeness Ratio (ECR): The Falsifiability Metric of MPAM↔MIEVM Duality

Document ID: ERES-ECR-2026-001

Author: Joseph A. Sprute (JAS), ERES Maestro

Affiliation: ERES Institute for New Age Cybernetics, Bella Vista, Arkansas

ORCID: 0000-0001-9946-3221

Date: 21 May 2026

Version: v1.4 (MIEVM Round 3 — Claude / ChatGPT / DeepSeek / USE.ai / Gemini)

License: CARE Commons Attribution License v2.1 (CCAL v2.1)

Series: ERES Trilogy support corpus — Volume Zero (Infomediary) instrumentation layer

Executive Summary

This summary is acronym-light and intended for first-time readers. Full canonical terminology is used from §1 onward; Appendix A contains the complete reference.

What ECR is. The Erasure Completeness Ratio (ECR) is a number in $[0, 1]$ that answers a single question for any governance system: “when this system says something is gone — a permission, a record, an attestation, an identity — is it actually gone, everywhere it ever propagated?” ECR is the ratio of confirmed-destroyed instances to the total instances tracked through the system’s propagation graph.

Why ECR matters. Most governance frameworks make erasure claims that cannot be checked. A regulator, an auditor, or a third party has no instrument by which to verify or falsify the claim. ECR is that instrument: it converts an unfalsifiable assertion (“the data is gone”) into a measurable, externally re-derivable quantity (“ $ECR(W) = 0.918$ over window W , with the propagation graph and revocation list published”).

How ECR cannot be gamed. Three properties prevent strategic inflation: (1) the denominator is estimated by stratified discovery sampling and reported at the upper bound of a 95% confidence interval, biasing ECR downward (§4.7); (2) a node that conceals its instances is treated as residual leakage, not as absent (§4.7); (3) four independent computational instruments must agree pairwise within 0.07 or the result is reported as VEILED (§6.4).

How ECR fails honestly. ECR cannot reach 1.0 in finite time with finite resources — the thermodynamic and propagation constraints forbid it. The whitepaper makes this explicit: ECR

approaches 1 asymptotically, and the residual ($1 - \text{ECR}$) must be contained, not denied. A system that promises $\text{ECR} = 1.0$ is dishonest by construction.

What you can verify yourself. If a deployment publishes its propagation graph, its revocation list, its verification keys, and its node roster, you can compute your own ECR' and compare it to the system's claim. The whitepaper provides a worked example (§11) where every number is independently re-derivable and a five-phase open-tooling roadmap through 2027 (§12).

What this whitepaper is, and is not. It is a specification with worked example and roadmap. It is not yet a completed pilot deployment; §12.6 names the external dependencies (independent auditor, jurisdictional partners, medical-systems integrator) honestly rather than concealing them.

A note on terminology. Where this paper uses the term "Aura" in the Aura Resonance Index (ARI), it refers specifically to the measurable bio-electromagnetic field signature established in the ERES Bio-Electric Signature work (ResearchGate publication 373331163, 2023), comprising heart-rate variability, EEG-derived coherence, and thermal-emission patterns aggregated under the FAVORS biometric layer. The term "Aura" is canonical to the ERES corpus; the underlying measurement is conventional bio-instrumentation. External publications and standards submissions may use the alternative external label "Bio-Electric Coherence Index (BECI)" as a transparency aid; ARI remains canonical inside the ERES corpus. See §8 for the full operational reduction.

What carries the legal load. The Daubert/Frye-admissible instrument is ECR computed by the deterministic reference implementation (§6.6, Phase 1 deliverable). The MIEVM-LLM ensemble (Claude, ChatGPT, DeepSeek, USE.ai, Gemini) is a governance-commentary instrument used for peer-review and pairwise convergence checking; it is not load-bearing for legal evidence. BERA / Resonance metrics inform governance tuning and are not offered as legal evidence.

Abstract

The Erasure Completeness Ratio (ECR) is the empirical falsifiability metric of ERES governance, first formalized in the March 2026 SPT specification within the SCALULAR ENGINE document. Canonically, $\text{ECR} = N_{\text{confirmed_destroyed}} / N_{\text{total_propagation_graph}}$: a value in $[0, 1]$ that measures, against a CBGMODD-tracked propagation graph, the fraction of instances of a target object (attestations, permissions, claims, data records, ratified states) that have been verifiably destroyed across every node on which they ever propagated. ECR is bounded above by the Triune Math key $\text{REAL} = (E \cdot M \cdot R) / (T \cdot S)$, so $\text{ECR} \rightarrow 1$ asymptotically and the residual ($1 - \text{ECR}$) must be contained rather than denied. ECR is the closure metric of the MPAM↔MIEVM Duality — the Melting Pot Assimilation Method (affirmative ledger) and the Multi-Instrument Ensemble Validation Method (operational runtime) — with ECR's asymptotic-yet-falsifiable character converting ERES from an abstract narrative system into a

governable, empirically testable control law over a one-thousand-year horizon. This whitepaper restates the canonical formula, anchors it within $C = R \times P \div M$ and the REAL bound, specifies its instrumentation through EAAP and EAAP-CRL under ML-DSA/Falcon post-quantum signatures, fixes the SCALULAR pillar thresholds (HEALTH ≥ 0.99 , PROTECTION ≥ 0.97), and demonstrates the seven-vector application across the RATE dimensions $R_1 \dots R_7$.

1. Background and Provocation

ERES has been challenged — appropriately and repeatedly — on whether its symbolic apparatus is empirically validatable, or whether the equations, backronyms, and lock-boxes constitute an abstract narrative system indistinguishable from rhetoric. The most recent and most precise form of this challenge, raised in May 2026 by an Execution Governance Strategist working on runtime oversight for autonomous AI systems, is reproduced in paraphrase:

The issue is less the symbolic equation itself and more whether the underlying governance logic can be empirically validated. Without definitional clarity around the variables and control structure, these models become difficult to distinguish from abstract narrative systems.

This challenge is legitimate. It is also, within ERES, already answered — but the answer requires naming the specific instrument that closes the loop, rather than restating the equation. That instrument is ECR. The remainder of this document specifies it.

2. The Canonical Root

ECR cannot be defined in isolation. It is the falsifiability metric of a specific control law, and that control law must be stated first, in its canonical form, with no variable migration.

2.1 Triune Math Key #1

$$C = R \times P \div M$$

Cybernetics = Resource $\times$ Purpose $\div$ Method

The variables are fixed. Drift in any of them — substituting Consequence for Cybernetics, Permissions for Purpose, Mitigation for Method — produces a different equation wearing the same letters, and is itself a failure mode (repeatable drift) that ECR is partly designed to detect across instruments.

2.2 The Three Governing Principles

Every ECR computation is subordinated to:

1. Don't hurt yourself.
2. Don't hurt others.
3. Build for generations to come.

An ECR of 1.0 obtained at the cost of violating any of these three is, by ERES doctrine, a null result.

2.3 Administration

$C = R \times P \div M$ is administered by the User-GROUP Community of Interest (COI) operating through GAIA. GAIA carries a deliberate dual reference: as institutional backronym it names the Global Actuary Investor Authority — the federated governance body responsible for implementing ERES standards and overseeing planetary resource allocation as the actuary of the planet; as substrate referent it names the biosphere itself — the planet as audit ground on which the institutional GAIA must report. HowWay is its operational handle. ECR is the number that tells the COI whether HowWay is actually doing what HowWay claims to do, with the actuarial GAIA reporting to and constrained by the biospheric GAIA.

3. The MPAM↔MIEVM Duality

Every ERES governance act has two faces:

MPAM (affirmative side)

Melting Pot Assimilation Method — the formal modular protocol by which mathematical frameworks, attestations, permissions, claims, and ratified states are admitted into the ERES corpus only when needed, as needed, and where needed. MPAM establishes the layered architecture, trigger-based adoption rules, containment tiers, and provenance-locking requirements that constitute the corpus of things ERES has positively assimilated. In the governance frame addressed by this whitepaper, MPAM is the affirmative ledger of what has been said yes to.

MIEVM (validation side)

Multi-Instrument Ensemble Validation Method — the cross-instrument check, currently realized through a short-list ensemble (Claude, Grok, DeepSeek, ChatGPT) culled by bandwidth constraint from a larger original set. MIEVM is the operational runtime of MPAM: it tests whether material admitted under MPAM holds up under independent re-derivation by multiple instruments. Where MPAM specifies what may be assimilated, MIEVM determines whether the assimilation actually converged.

The duality claim

MPAM is only as trustworthy as MIEVM can re-derive it, and MIEVM is only as trustworthy as the erasure side of MPAM is complete. A framework, attestation, or claim assimilated under

MPAM that cannot be cleanly de-assimilated when its triggers fail is a framework that was never really under control. The duality is therefore not a redundancy — it is a closure requirement, with de-assimilation (Tier 5 Archive) as the structural mirror of assimilation. ECR is the metric of that closure.

4. Erasure Completeness Ratio — Formal Definition

4.1 The canonical ratio

The Erasure Completeness Ratio was first formalized in the March 2026 SPT specification within the SCALULAR ENGINE document, in the context of CBGMODD's reframed role as erasure-completeness auditor rather than binary-certificate issuer. The canonical formula is:

$$\text{ECR} = \text{N_confirmed_destroyed} \div \text{N_total_propagation_graph}$$

Where `N_confirmed_destroyed` is the number of confirmed-eliminated instances of the target object (including copies, caches, backups, derivative products, signed attestations, granted permissions, and ratified claim states) and `N_total_propagation_graph` is the total number of instances identified by CBGMODD's provenance tracking across the object's full lifecycle. ECR is a value in $[0, 1]$.

4.2 The REAL bound

ECR is structurally bounded by the third Triune Math key, $\text{REAL} = (E \cdot M \cdot R) / (T \cdot S)$:

$$\text{ECR_max} = f(E, M, R, T, S)$$

Erasure effectiveness is bounded by the Energy and Matter resources that can be applied, modulated by the Resonance of the system's coordination capacity, across the Time and Space over which the object has propagated. This is the same thermodynamic frontier that bounds complete recycling of matter — there is always residual signal (magnetic remnance, unenumerated cache fragments, derivative hashes, neural traces in human readers, propagated copies in surfaces not yet discovered by CBGMODD). The honest engineering position is therefore not “ECR can reach 1” but rather “ECR approaches 1 asymptotically and the residual is what we contain.” When the marginal cost of pursuing the next data fragment exceeds the REAL value of the system, the practical limit of erasure has been reached for that window.

4.3 Interpretation

- **ECR → 1.0** — asymptotic ceiling; full closure within the resource envelope; residual is bounded by REAL.
- **ECR < 1.0** — residual entries persist beyond the REAL frontier: stale signatures still verifying, revoked permissions still being honored, retracted claims still propagating

downstream. The complement ($1 - \text{ECR}$) is the leakage fraction; it must be contained, not denied.

- **ECR = VEILED** — reserved state for windows in which the propagation graph itself is unresolvable (the denominator cannot be enumerated, typically because CBGMODD has not yet mapped a surface that may hold residual copies). VEILED replaces the earlier $\perp$ (bottom) notation; it is not zero, not undefined, and not failure — it is the explicit acknowledgement that the question is not currently answerable, which itself is information.

4.4 The seven RATE dimensions

ECR is not a single scalar in production deployments. Following the canonical relation $(\text{CBGMODD} \times \text{FAVORS}) + \text{BERA} = \text{RATE}$, which yields seven RATE dimensions R_1 through R_7 , a production ECR is reported as the vector $\langle \text{ECR}_1, \text{ECR}_2, \dots, \text{ECR}_7 \rangle$ (one ratio per RATE dimension), and the convergent mean across the seven is what is compared to the MIEVM GO threshold (≥ 8.5). A scalar reported without the seven-vector behind it is presumptively VEILED. This is the whitepaper's synthesis: the canonical ECR formula applied dimension-wise across RATE.

4.5 Why this is empirically testable

Each term in the canonical formula is operationally defined and externally measurable:

- $N_{\text{total_propagation_graph}}$ is enumerated by CBGMODD's provenance tracking against the COI charter and the EAAP-CRL revocation list — all public, all signed, all auditable.
- $N_{\text{confirmed_destroyed}}$ is determined by polling every node in the propagation graph (registries, mirrors, cached attestations, replica nodes, derivative products) and confirming that the target instance is no longer present or that signatures over it no longer verify.
- The numerator is a subset of the denominator by construction; both are finite within the operational window.

A narrative claim that “the system retires obsolete permissions” cannot be falsified. ECR can be falsified, in principle, by any third party with access to the EAAP-CRL, the CBGMODD propagation graph, and the downstream surfaces; if their independent computation diverges from the ERES-reported ECR, the discrepancy is itself a measurable error term. This is the property that converts the system from narrative to control law.

4.6 SCALULAR pillar thresholds

ECR floors are not uniform across the four SCALULAR pillars; each pillar carries the threshold appropriate to the harm domain it governs:

- **HEALTH (SSHP):** $\text{ECR} \geq 0.99$ — patient data, medical records, and bio-identifying material require the highest erasure floor.

- **PROTECTION (SSPS):** $ECR \geq 0.97$, adjusted downward for inferential-reconstruction risk — whistleblower identity, witness protection, and fear-of-retribution elimination cases.
- **LAW (SSLA) and SKILLS_TRADE (SSST):** thresholds set by the adopting jurisdiction within the COI charter, never lower than 0.90 for any data class touching personally-identifying material.

These thresholds anchor the SCALULAR ENGINE's fear-of-retribution elimination requirement: protecting a disclosure is only credible if the identity of the discloser can be measurably erased to the pillar-appropriate floor.

4.7 Denominator discovery and non-cooperative nodes

The canonical formula presupposes that $N_{\text{total_propagation_graph}}$ is enumerable. In practice, CBGMODD's provenance tracking may be incomplete — some surfaces are partition-resistant (darknet caches, air-gapped replicas, jurisdictionally hostile mirrors), and some nodes may decline to disclose their instance set. v1.2 introduces an explicit discovery protocol and a non-cooperative-node clause closing this gap.

Discovery protocol

CBGMODD performs a two-stage enumeration: (a) declared nodes from the COI charter and EAAP-CRL routing record; (b) sampled discovery against the wider network using a stratified sampling strategy over jurisdictional, technological, and temporal strata. The estimator $\hat{N}_{\text{total}}$ is reported with a 95% confidence interval; $N_{\text{total_propagation_graph}}$ in the published ECR uses the upper bound of that interval. This makes the denominator conservatively large, biasing the reported ECR downward rather than upward — the honest direction of error.

4.7.1 Confidence-bound formalization

v1.3 formalizes the discovery-sampling confidence bound. Let N_{declared} denote the count of declared nodes (stage a). Stage (b) draws a stratified sample of size n across S strata indexed by jurisdiction, technology surface, and temporal cohort, with stratum capture rate $\hat{p}_s$ estimated within each stratum. The point estimate is $\hat{N}_{\text{total}} = N_{\text{declared}} + \sum_s \hat{N}_{\text{undeclared},s}$, where $\hat{N}_{\text{undeclared},s}$ is the Horvitz–Thompson estimator within stratum s . The 95% upper confidence bound is computed as $\hat{N}_{\text{total}} + 1.96 \cdot \sigma(\hat{N}_{\text{total}})$, with σ the standard error pooled across strata. The published denominator uses this upper bound, with N_{declared} , n , S , the per-stratum $\hat{p}_s$, and σ themselves published as audit fields. A third party with access to these fields can independently re-derive the upper bound and recompute ECR'; any discrepancy beyond rounding is a Level-2 finding under §6.5. This protocol does not eliminate denominator uncertainty — truly dark surfaces (human-memorized derivatives, air-gapped replicas not in any declared stratum) remain unenumerable — but it bounds the uncertainty within published confidence and makes the bound itself externally checkable.

Non-cooperative-node clause

A node that refuses to disclose its instance set, or that cannot be polled within the operational window, is treated by default as a residual contribution to $(1 - \text{ECR})$ until validated by a COI-ordered forensic audit. Persistent refusal across two consecutive windows promotes the affected object's ECR to VEILED with a non-cooperative-node flag, and triggers escalation under §4.8. This means an adversary cannot improve the apparent ECR by hiding propagation; concealment is mathematically equivalent to admitted leakage.

4.8 VEILED governance

VEILED is not a stable resting state. v1.2 specifies entry, residency, and exit criteria so that VEILED cannot mask uncertainty indefinitely.

- **Entry:** an object enters VEILED when (a) CBGMODD cannot enumerate at least one surface in its propagation graph, (b) a node returns non-cooperative for two consecutive windows, or (c) MIEVM pairwise disagreement exceeds the §6.4 bound.
- **Residency:** default time-box of 30 days, extendable once by COI vote to 90 days with published reasoning. Residency time is itself a published audit field.
- **Exit (resolution):** the propagation graph is closed (numeric ECR restored), or remediation completes against the non-cooperative node, or MIEVM re-converges.
- **Exit (failure):** if residency expires without resolution, the object is automatically downgraded — its ECR floor is treated as the pessimistic lower bound (assume all unresolved instances are residual), the affected SCALULAR pillar service is held in CONTAINED state, and a §6.5 drift alarm is raised against the COI.

VEILED therefore has consequences. It is information, not refuge.

5. Instrumentation

ECR is not a calculation performed in isolation; it sits on a stack of named ERES instruments, each of which is itself externally inspectable. The table below maps each sub-component of the ECR computation to the instrument that produces it and to the mode by which a third party can falsify the reported value.

ECR sub-component	ERES instrument	Falsifiability mode
N_total_propagation_graph enumeration	CBGMODD provenance tracking + EAAP-CRL revocation list (post-quantum signed)	Third-party verifies CRL signature, recomputes the propagation graph independently

ECR sub-component	ERES instrument	Falsifiability mode
N_confirmed_destroyed verification	Downstream node polling under ML-DSA / Falcon signatures	Independent re-verification of instance absence or signature non-validity
Subject anchoring	FAVORS six-factor biometric (Fingerprint, Aura, Voice, ODOR, Retina, Signature)	Per-subject erasure traceable to a unique embodied claimant
Erasure-completeness audit (not binary certificate)	CBGMODD acting as completeness auditor; outputs ECR with timestamp	Audit log is signed and re-derivable by any COI member
Cross-instrument check	MIEVM ensemble (Claude, Grok, DeepSeek, ChatGPT) at convergent mean ≥ 8.5	Divergent ECR across instruments → repeatable-drift flag
Tuning signal	BERA Resonance Metrics: ARI (Aura Resonance Index), ERI (Emission Resonance Index), RHC (Resonant Harmony Cycle), RCI (Resonant Continuity Index)	ECR feeds Meritcoin under Proof-of-Resonance (“tuning, not mining”)
Thermodynamic ceiling	$REAL = (E \cdot M \cdot R) / (T \cdot S)$ — Triune Math Key #3	$ECR_{max} = f(E, M, R, T, S)$; the residual $(1 - ECR)$ is contained, not denied
Unresolvable state	VEILED (supersedes $\perp$)	Explicit non-answer; itself audit information

Each row is independently auditable. The full ECR cannot be fabricated without simultaneously fabricating every instrument in the right-hand column — which is precisely the property that converts the system from narrative to control law.

6. ECR as the Falsifiability Hook

The provocation in Section 1 reduces, formally, to: “provide an operationally defined quantity Q such that an external party can compute Q' from the same inputs and compare Q' to ERES-reported Q .” ECR is that quantity.

6.1 The independent re-derivation test

For any operational window, an external auditor with access to:

- the CBGMODD propagation graph for the target object (the denominator $N_{\text{total_propagation_graph}}$),
- the published EAAP-CRL for the window,
- the set of downstream nodes declared in the COI charter,
- the post-quantum verification keys (ML-DSA / Falcon public keys),

can compute their own $ECR' = N'_{\text{confirmed_destroyed}} / N'_{\text{total_propagation_graph}}$ without any privileged cooperation from ERES. The discrepancy $\Delta = |ECR - ECR'|$ is itself a measurable error term, and a non-zero Δ is a falsification event — either an ERES reporting error, an auditor methodology error, or a substrate compromise. All three are actionable, and the partitioning of Δ across the seven RATE dimensions usually identifies which.

6.2 The repeatable-drift detector

A particularly important class of failure that ECR detects: repeatable drift. If across multiple windows $W_1, W_2, \dots, W_n$ the seven-vector $\langle ECR_1, \dots, ECR_7 \rangle$ trends monotonically downward in any dimension without an explanatory event, the system is leaking erasure in a structured way — quietly accumulating residual permissions or stale attestations. This is the AI-determinism failure mode (deterministic but drifting) reframed in governance terms, and it is distinct from ordinary random error. ECR makes it visible.

6.3 The narrative-system exclusion

A pure narrative system cannot produce an ECR because it has no enumerable propagation graph. The denominator does not exist. Any system that publishes an ECR (with the canonical formula, the seven-vector, and the CBGMODD propagation provenance) has, by that publication alone, exited the narrative-system category. ECR is therefore the bright line.

6.4 MIEVM pairwise disagreement bound

Convergent-mean reporting can mask significant pairwise disagreement among instruments. v1.2 specifies a pairwise bound to close this gap:

$$\max_i, j |ECR_i - ECR_j| \leq 0.07$$

If the maximum pairwise divergence across the MIEVM ensemble exceeds 0.07 (seven percentage points), the reported ECR is automatically promoted to VEILED, the full divergence matrix is published, and the responsible instruments are flagged for re-derivation. This rule prevents two instruments from agreeing closely while two others disagree sharply — a configuration that yields an acceptable mean but conceals genuine epistemic disagreement. The 0.07 figure was selected to align with the SCALAR PROTECTION pillar floor distance ($1.0 - 0.97 = 0.03$) plus the HEALTH-pillar tolerance (0.01) plus a methodological safety margin (0.03); future MIEVM rounds may tighten or relax this bound with explicit COI ratification.

6.5 Drift alarm control limits

Beyond single-window verification, v1.2 specifies pre-registered control limits for the ECR time series, so that drift is detected algorithmically rather than discretionally:

- **Level-1 advisory:** any single-window ECR_i drops below the SCALULAR pillar floor; service for that object class is held in WATCH state pending the next window.
- **Level-2 warning:** monotone downward trend in any ECR_i across three consecutive windows without an explanatory event in the EAAP-CRL log; mandatory COI review.
- **Level-3 alarm:** the convergent mean across all seven RATE dimensions drops below the MIEVM GO threshold (≥ 8.5 normalized), or any $\Delta = |ECR - ECR'|$ exceeds 0.05 against an external auditor; mandatory CONTAINED-state hold on the affected SCALULAR pillar service until remediated.
- **Level-4 escalation:** two consecutive Level-3 alarms within the same operational period escalate to the 1000-year Equity Covenant trustees with public publication of the divergence record.

Control limits are pre-registered, not selected post hoc. This is the property that converts ECR drift from a narrative observation into an actionable signal.

6.6 Deterministic Reference Implementation and the LLM-Ensemble Boundary

v1.4 introduces an explicit separation between two computational instruments operating in this specification, both required, neither substitutable for the other.

Deterministic reference implementation (load-bearing)

The legally load-bearing instrument is the deterministic reference implementation: an open-source ECR computation library that takes as input the published audit fields (CBGMODD propagation-graph snapshot, EAAP-CRL excerpt, ML-DSA / Falcon public verification keys, COI charter, discovery-sampling audit fields per §4.7.1) and produces a bit-identical output across all conformant implementations. Two independent implementations of the reference library must produce the same ECR seven-vector $\langle ECR_1 \dots ECR_7 \rangle$ to within the published precision, or one of them has a defect. This deterministic property is what satisfies the Daubert testability and known-error-rate factors. The deterministic library is the Phase 1 deliverable per §12.1 (2026-Q3 target).

MIEVM-LLM ensemble (governance-commentary instrument, not load-bearing)

The MIEVM-LLM ensemble (Claude, ChatGPT, DeepSeek, USE.ai, Gemini; Grok when available) is a governance-commentary instrument that performs peer-review, gap-identification, doctrinal cross-checking, and pairwise convergence on the framework itself — not on the ECR numerical computation. The pairwise disagreement bound of 0.07 (§6.4) applies to ensemble *commentary*, including ratings of the specification's coherence, completeness, and falsifiability. It does not apply to numerical ECR values, which are produced by the deterministic library and are bit-identical by construction.

Why both are needed

The deterministic library produces a number; it does not assess whether the number is meaningful in context, whether the propagation graph is complete in the relevant sense, whether the threshold is appropriate for the harm domain, or whether the methodology has drifted from canon. Those are governance-commentary questions. The MIEVM-LLM ensemble handles them. The legal evidence is the deterministic output plus its audit fields; the governance assurance is the ensemble commentary plus its pairwise bound.

Daubert / Frye implications

Under *Daubert v. Merrell Dow Pharmaceuticals*, 509 U.S. 579 (1993), expert evidence requires testability, peer review, known error rates, standards control, and general acceptance. The deterministic reference implementation satisfies factors 1–4 by construction: testable by independent re-implementation, peer-reviewable through open-source release, with known error rates bounded by floating-point precision and audit-field accuracy, and methodology fixed in advance via the canonical formula and pre-registered control limits. Factor 5 (general acceptance) is the proper subject of standards engagement per §12 and the Legal v1.2 administrative pathways. The MIEVM-LLM ensemble does not need to satisfy Daubert because it is not offered as evidence; it is offered as peer-review process.

6.7 Adversarial Model and Resistance Analysis

v1.4 introduces an explicit threat model. ECR's defenses are described above per attack class; what follows is the consolidated adversarial-model statement that any cryptographic-protocol paper conventionally includes and that serious reviewers should be able to evaluate.

Attack class A: Concealment (non-cooperative node)

An adversary refuses to disclose its instance set. Per §4.7, the affected node's contribution is treated as residual leakage ($1 - \text{ECR}$) until validated by a COI-ordered forensic audit. Persistent refusal over two consecutive windows promotes the object to VEILED with the non-cooperative-node flag (§4.8). Adversary outcome: cannot inflate ECR by hiding. The ratio degrades honestly to reflect the hidden surface.

Attack class B: Sybil enumeration

An adversary submits false propagation-graph nodes to inflate the denominator and artificially reduce ECR (e.g., to manufacture a VEILED state and trigger procedural delay). Defense: the COI charter constrains declared-node admission under MPAM. Stratified discovery sampling per §4.7.1 validates the declared set against an independently drawn stratified sample; large discrepancies trigger Level-2 warning under §6.5. Adversary outcome: false nodes are detectable as a divergence between declared and sampled enumeration.

Attack class C: CRL forgery

An adversary forges revocation entries pre-key-compromise. Defense: EAAP-CRL signatures are ML-DSA (for general attestation) and Falcon (for size-constrained CRL entries), with security against quantum and classical adversaries per the FIPS 204 and NIST PQC standards. The Phase 3 absence-proof primitives (§12.3) add a transparency-log layer (RFC 6962-style Merkle commitments) so that any CRL state is reconstructable from the log even if key compromise occurs later. Adversary outcome: forgery requires breaking both the post-quantum signature and the transparency log.

Attack class D: Air-gapped retention (the dark-surface limit)

An adversary maintains an offline replica that never appears in any declared or discoverable stratum. Defense: none, by construction. ECR cannot measure what cannot be enumerated. This is the permanent technical limit of the framework, and v1.4 acknowledges it directly. The legal implication is treated in Legal v1.2 §2.9: ECR provides preponderance-level evidence of erasure within the enumerable surface but does not claim clear-and-convincing or beyond-reasonable-doubt evidence for the dark surface. Parties offering ECR as legal evidence must state this scope limit on the record.

Attack class E: Coerced node

An adversary compels a node operator to falsify a polling response (e.g., claim destruction when retention persists). Defense: ECR Evidence Bundle auditor declaration under penalty of perjury (28 U.S.C. §1746 in the U.S. federal context, or analogous foreign instrument); deterministic recomputation by an independent auditor surfaces the discrepancy as $\Delta = |\text{ECR} - \text{ECR}'|$. Persistent unexplained Δ above 0.05 triggers Level-3 alarm per §6.5. Adversary outcome: coercion is detectable but not prevented; the framework converts coercion from concealment into a measurable governance failure.

Attack class F: Governance capture

An adversary takes over the COI itself, controlling charter, declared nodes, and audit access. Defense: the 1000-year Equity Covenant trustees provide Level-4 escalation independent of the operational COI. Independent re-derivation by parties outside the captured COI — using only the public audit fields, deterministic library, and EAAP-CRL transparency log — remains possible. Adversary outcome: capture is detectable through external re-derivation divergence; the framework does not prevent capture but prevents capture from being concealed.

What this analysis is, and is not

This is a governance threat model, not a cryptographic security proof. It enumerates attack classes and their treatment under the framework's defenses; it does not claim that all defenses are mathematically complete. Two defenses (post-quantum signatures, transparency-log absence proofs) are cryptographic in the strict sense. The others are governance instruments — procedural, evidentiary, escalation-based — and provide reasonable assurance rather than mathematical certainty. The honest legal framing of this distinction is the subject of Legal v1.2 §2.9.

7. ECR Under Post-Quantum Conditions

ECR is meaningful only as long as the signatures it depends on remain meaningful. EAAP migrated to ML-DSA and Falcon in early 2026 specifically to ensure that erasure semantics remain falsifiable on a one-thousand-year horizon (per the Equity Covenant, ERES-COVENANT-2026-001).

Under classical signature schemes, a sufficiently advanced adversary could forge revocation entries or, conversely, forge proofs that revoked objects remain valid. Either attack collapses ECR — the metric returns plausible-looking numbers while the underlying erasure surface is being silently rewritten. ML-DSA (for general attestation) and Falcon (for size-constrained CRL entries) push the falsifiability horizon out beyond foreseeable cryptographic advances, including those leveraging quantum computation.

This is not theoretical hardening. It is the load-bearing requirement for ECR to remain a 1000-year metric rather than a 25-year metric.

8. ECR and Proof-of-Resonance

Meritcoin operates under Proof-of-Resonance, not Proof-of-Work. The distinction matters for ECR specifically:

Proof-of-Work has no native erasure semantics

Once a block is buried under sufficient subsequent work, it is buried. There is no operational notion of “retiring” that block — the chain treats it as permanent. An ECR-equivalent computed against a PoW substrate would always be near zero in the denominator-defined sense (almost nothing is ever required to be erased) and therefore vacuous.

Proof-of-Resonance is continuously revisable

Resonance state is a tuning, not a settlement. The four BERA Resonance Metrics — ARI (Aura Resonance Index, the personal/biometric coherence field), ERI (Emission Resonance Index, the environmental/ecological output measurement), RHC (Resonant Harmony Cycle, the cyclical feedback and temporal rhythm), and RCI (Resonant Continuity Index, the cross-scale persistence metric, formalized in collaboration with Jimmy D. Butzbach’s Continuity Theory as $RCI = P_Q_norm \times ARI_sys \times VibConst$) — are re-evaluated at every Meritcoin update window. A prior tuning that no longer holds is not buried; it is retired, and the retirement is exactly what ECR measures. “It’s not mining — it’s tuning” is therefore not a slogan but a structural claim about the erasure surface: every Meritcoin update presupposes that the erasure of superseded resonance states has been audited by ECR.

Operational note on ARI terminology (extended in v1.4). MIEVM reviewers across Rounds 2 and 3 flagged that “Aura” risks pseudoscience classification by readers unfamiliar with the ERES corpus. Within ERES, “Aura” denotes the measurable bio-electromagnetic field signature established by the ERES Bio-Electric Signature (BES) work (ResearchGate publication 373331163, 2023), comprising HRV, EEG-derived coherence, and thermal-emission patterns aggregated under the FAVORS biometric layer. The operational reduction: ARI is the coherence measurement over the BES feature set, with reference instrumentation specified in the BES paper. Reviewers wishing to evaluate ARI empirically should treat “Aura” as a label for that measurement, not as a metaphysical claim.

External alternative label (v1.4 addition). External publications, standards submissions to NIST / IEEE / ISO/IEC, and regulatory filings may use the alternative external label “Bio-Electric Coherence Index (BECI)” as a transparency aid for non-ERES audiences. ARI remains canonical inside the ERES corpus; BECI is offered as a one-to-one alias for boundary-spanning communication. This dual-naming pattern is the standard approach for canonical technical terms whose connotations differ outside their native domain.

Legal-load disclaimer (v1.4 addition). BERA / Resonance metrics including ARI, ERI, RHC, and RCI are governance-tuning instruments. They are not offered as legal evidence under Daubert or Frye. The legally load-bearing instrument is ECR computed by the deterministic reference implementation (§6.6). Parties wishing to adopt ECR can do so without taking any position on the BERA metrics; the metrics inform Meritcoin tuning and ERES corpus governance, not legal admissibility.

ECR is the metric that audits whether the tuning system is honoring its own retirement obligations. Without ECR, Proof-of-Resonance would have the same auditability gap as Proof-of-Work; with ECR, it has the closure property that PoW lacks by construction.

9. The One-Thousand-Year Horizon

The ERES Equity Covenant (ERES-COVENANT-2026-001) is a one-thousand-year governance instrument. ECR is the metric that makes that horizon empirically meaningful rather than aspirational. Three properties together produce this:

4. Each ECR computation is bounded in a specific window W ; the horizon is composed of contiguous windows, each independently verifiable.
5. The instrumentation (EAAP-CRL under ML-DSA/Falcon, FAVORS biometric anchoring, MIEVM cross-instrument check) is designed to remain falsifiable across substrate generations — cryptographic, computational, and institutional.
6. VEILED windows are explicitly preserved as VEILED, not retroactively coerced into numerical values. A 1000-year record will contain VEILED entries; their preservation is itself part of the audit substrate.

The horizon, in other words, is not maintained by claiming completeness; it is maintained by maintaining the apparatus that makes incompleteness visible.

10. MIEVM Validation Protocol for This Whitepaper

In accordance with the ERES Force-Multiplier Method v1.5, this document is subject to MIEVM cross-instrument validation prior to canonical deposit. The GO threshold is a convergent mean ≥ 8.5 across the four instruments (Claude, Grok, DeepSeek, ChatGPT), with attention to:

- Canonical preservation: variables in $C = R \times P \div M$ must not migrate.
- Subject-First Nomenclature: ECR is named for what it measures (erasure completeness), not for what it serves.
- Falsifiability: every numerical claim must be reproducible by an external party from declared inputs.
- Drift detection: the seven-vector $\langle ECR_1 \dots ECR_7 \rangle$ across windows must be preserved, not collapsed to scalar.

Convergence below threshold triggers revision. VEILED is permitted as an instrument response and does not by itself block deposit, provided the VEILED reasoning is itself reported.

11. Worked Example — Whistleblower Disclosure Under SCALULAR PROTECTION

This section makes the canonical formula concrete with a single end-to-end example. All numbers are illustrative; the structure is canonical.

11.1 Scenario

Operational window: $W = 2026\text{-}Q2$, length 90 days. Target object: personally-identifying information (PII) associated with whistleblower disclosure ID-WB-2026-042, governed under the SCALULAR PROTECTION pillar (SSPS, floor $ECR \geq 0.97$ with inferential-reconstruction adjustment). At the start of W , the COI has issued an EAAP-CRL entry revoking all attestations referencing the discloser's identity; CBGMODD begins propagation-graph enumeration.

11.2 Propagation-graph enumeration

CBGMODD enumerates 47 declared nodes carrying the target object:

- 12 EAAP attestation surfaces (primary credentials)
- 8 cached attestation mirrors (distributed across jurisdictions)
- 14 downstream replica nodes (jurisdictional partner registries)

- 7 derivative-product surfaces (audit logs containing PII references)
- 6 backup vaults (operated by the COI itself)

Stratified discovery sampling (per §4.7) yields an additional 2 nodes not in the declared list — unenumerated cached mirrors at non-COI infrastructure providers. CBGMODD updates the denominator: $N_{\text{total_propagation_graph}} = 49$ (upper bound of the 95% confidence interval).

11.3 Erasure cascade and ECR computation

Over W , revocation completes across 45 nodes. Two backup vaults are on a scheduled erasure cycle that extends into $W+1$, and the two newly-discovered unenumerated mirrors return non-cooperative responses to polling. Under the §4.7 clause, the two non-cooperative mirrors are treated as residual contributions to $(1 - \text{ECR})$.

$$\text{ECR}(W) = 45 / 49 = 0.918$$

This is below the PROTECTION floor of 0.97. The PROTECTION pillar service for ID-WB-2026-042 is held in CONTAINED state. The two non-cooperative mirrors are flagged under §4.7; remediation is initiated through jurisdictional channels.

11.4 The independent re-derivation

An external auditor with access to the CBGMODD propagation graph, the EAAP-CRL, and the ML-DSA / Falcon verification keys recomputes against the same surfaces and discovers one additional jurisdictional replica node that CBGMODD had not enumerated. Their computation:

$$\text{ECR}'(W) = 45 / 50 = 0.900$$

$$\Delta = |0.918 - 0.900| = 0.018$$

Δ is below the §6.5 Level-3 alarm threshold of 0.05, but it is a measurable error term. The auditor publishes their finding; CBGMODD ingests the additional node into its propagation graph; the canonical ECR is updated to $45/50 = 0.900$. The PROTECTION pillar remains in CONTAINED state.

11.5 The VEILED case

One of the two non-cooperative mirrors is located in a partition-resistant darknet replica. The set of instances at that mirror is not enumerable by external means. The sub-graph terminating at that node is reported as VEILED under §4.8 with reason “non-enumerable mirror at node ND-4127.” Default residency: 30 days. COI extends to 90 days by published vote, citing diplomatic remediation in progress. If unresolved at day 90, the object’s ECR is automatically downgraded to the pessimistic lower bound (assume all unresolved instances are residual), and a §6.5 Level-3 drift alarm is raised against the COI.

11.6 The MIEVM cross-check

Four instruments independently compute ECR for ID-WB-2026-042 over W:

- Claude: 0.918
- ChatGPT: 0.910
- DeepSeek: 0.925
- USE.ai: 0.915

Convergent mean = 0.917. Max pairwise divergence = $|0.925 - 0.910| = 0.015$, which is below the §6.4 bound of 0.07. MIEVM converges; the reported ECR is canonical at 0.918 (Claude's value adopted as the publishing instrument), with the four-instrument record published as part of the audit trail. If max pairwise had exceeded 0.07, the ECR would have been promoted to VEILED with the full divergence matrix published.

11.7 What the example demonstrates

Every claim in this worked example is independently re-derivable from the published CBGMODD propagation graph, the EAAP-CRL for W, the ML-DSA / Falcon verification keys, and the COI charter. There is no step that depends on privileged ERES cooperation. The example is therefore itself an instance of the property the whitepaper claims for ECR: a third party can verify or falsify any number in this section without asking ERES for permission.

12. Reference Implementation Roadmap

Independent re-derivation requires public tooling. v1.2 commits to the following sequence of open releases through 2027, each milestone gated by MIEVM convergence and a published audit record.

12.1 Phase 1 — CBGMODD enumerator open release (2026-Q3)

Open-source release of the CBGMODD provenance-tracking enumerator with a synthetic test corpus, EAAP-CRL integration under ML-DSA, and a worked specification of the discovery-sampling estimator. Reproduction kit will allow third parties to enumerate propagation graphs for arbitrary attested objects against a public test substrate.

12.2 Phase 2 — Public ECR testnet (2026-Q4)

Public testnet hosting synthetic propagation graphs of escalating complexity (10^2 , 10^4 , and 10^6 nodes). Third parties run their own ECR' computation; the testnet publishes the Δ leaderboard. Convergence below $\Delta = 0.01$ across at least three independent implementations gates Phase 3.

12.3 Phase 3 — Absence-proof primitives (2027-Q1)

Integration of authenticated data-structure primitives (Merkle-tree commitments per RFC 6962, verifiable transparency logs, succinct non-membership proofs) so that `N_confirmed_destroyed`

claims are themselves cryptographically attestable rather than asserted. This is the load-bearing requirement for distinguishing “confirmed destroyed” from “asserted destroyed.”

12.4 Phase 4 — First production audit (2027-Q2)

First production audit against the ERES-MPAM-2026-001 corpus using the open tooling. ECR computed by three independent implementations; results published with full divergence matrix and Δ analysis. This is the empirical milestone that promotes the framework from specified to demonstrated.

12.5 Phase 5 — SCALULAR HEALTH pillar pilot (2027-Q3)

First operational SCALULAR pillar deployment, targeting HEALTH (SSHP) given its 0.99 floor and the relative maturity of medical-record erasure norms (HIPAA-aligned). Pilot includes published cost-risk model (per Appendix D) and continuous public ECR reporting.

12.6 Open dependencies

The roadmap depends on three things ERES cannot fund alone: an independent auditor willing to run the Phase 2 testnet, jurisdictional partners willing to host Phase 4 propagation surfaces, and a medical-systems integrator willing to host the Phase 5 pilot. v1.2 publishes the roadmap to make these dependencies visible; partner inquiries should reference ERES-ECR-2026-001 v1.2.

13. Conclusion

The Erasure Completeness Ratio is the specific empirical instrument that converts ERES governance from a symbolic apparatus into a falsifiable control law. It does this by making one question operationally answerable across any deployment window and any governance object class:

“When this system says something is gone, is it actually gone, everywhere?”

ECR returns a number in $[0, 1]$ computed against a CBGMDD-tracked propagation graph and bounded above by REAL, or it returns VEILED with documented entry/exit criteria and a 30-day default residency. Either return is information. Either return is auditable. Either return is independently re-derivable from the published surface.

v1.2 adds the operational closures the MIEVM ensemble identified as gaps in v1.1: a discovery-sampling estimator with conservative upper-bound denominators, a non-cooperative-node clause that converts concealment into admitted leakage, time-bounded VEILED governance with automatic downgrade on residency expiry, a pairwise MIEVM disagreement bound, pre-registered drift control limits across four escalation levels, a fully

worked example with concrete numbers and an explicit Δ divergence case, and a five-phase reference implementation roadmap through 2027.

$C = R \times P \div M$ holds. Administered by User-GROUP COI through GAIA — institutional Global Actuary Investor Authority reporting to biospheric audit ground — with HowWay as the operational handle. ECR is the metric by which HowWay demonstrates, window by window, that it is doing what it says it does. The instrument is specified, the falsifiability surface is published, and the implementation path is dated.

Appendix A — Canonical Cross-References

- ERES Triune Math Keys: (1) $C = R \times P \div M$ (Cybernetics = Resource $\times$ Purpose $\div$ Method); (2) $M \times E + C = R$; (3) $REAL = (E \cdot M \cdot R) / (T \cdot S)$.
- MPAM — Melting Pot Assimilation Method (formal modular protocol for admitting frameworks/attestations into ERES). Companion document: ERES-MPAM-2026-001.
- MIEVM — Multi-Instrument Ensemble Validation Method (operational runtime of MPAM; current short-list ensemble: Claude, Grok, DeepSeek, ChatGPT; GO threshold convergent mean ≥ 8.5).
- GAIA — dual reference: institutional backronym Global Actuary Investor Authority (federated governance body); substrate referent the biosphere itself (audit ground).
- ERES Equity Covenant — ERES-COVENANT-2026-001 (1000-year governance instrument).
- ERES Attestation and Authorization Protocol — EAAP v1.5, with EAAP-CRL under ML-DSA / Falcon.
- ERES Reliability Whitepaper — ERES-RELIABILITY-2026-001 (five AI failure vectors mapped to ERES instruments).
- ERES Force-Multiplier Method v1.5 — Subject-First Nomenclature; MIEVM GO threshold convergent mean ≥ 8.5 .
- ERES Infomediary (Volume Zero of the ERES Trilogy Architecture), consolidated v2.0.
- SCALULAR Engine Document — four pillars: HEALTH (SSHP, ECR floor ≥ 0.99), LAW (SSLA), PROTECTION (SSPS, ECR floor ≥ 0.97 with inferential-reconstruction adjustment), SKILLS_TRADE (SSST); Tier 1 = SSSC (Solid-State Smart-City Citizen). Original locus of the canonical ECR formula in the SPT specification (March 2026).
- UBIMIA Manual v1.1 (25 April 2026) — partial instruction set for the PlayNAC engine.
- FAVORS six-factor biometric: Fingerprint · Aura · Voice · ODOR · Retina · Signature.

- BERA Resonance Metrics: ARI (Aura Resonance Index) · ERI (Emission Resonance Index) · RHC (Resonant Harmony Cycle) · RCI (Resonant Continuity Index, with Butzbach formula $RCI = P_Q_norm \times ARI_sys \times VibConst$).
- CBGMODD: Citizen · Business · Government · Military · Ombudsman · Dignitary · Diplomat (seventh seat = Diplomat, not repeated Dignitary). CBGMODD acts as erasure-completeness auditor (not binary-certificate issuer) in the ECR context.
- $(CBGMODD \times FAVORS) + BERA = RATE$, yielding seven RATE dimensions R_1 through R_7 .
- VEILED — supersedes $\perp$ (bottom) as the semantically unresolvable RATE state.

Appendix B — Notation

- **ECR** = $N_confirmed_destroyed / N_total_propagation_graph$ — canonical Erasure Completeness Ratio (SPT specification, March 2026).
- **ECR_{max}** = $f(E, M, R, T, S)$ — REAL bound on ECR; the thermodynamic ceiling derived from Triune Math Key #3.
- **⟨ECR₁ ... ECR₇⟩** — ECR seven-vector across the RATE dimensions $R_1 \dots R_7$ (whitepaper synthesis applying the canonical formula dimension-wise).
- **1 – ECR** — leakage fraction; the residual that must be contained rather than denied.
- **VEILED** — explicit unresolvable state; replaces $\perp$.
- **Δ** = $|ECR - ECR'|$ — audit divergence between ERES-reported and externally-recomputed ECR.

Appendix C — Version History

v1.0 → v1.1 (canonical-correction pass)

- Corrected MPAM expansion: “Multi-Party Attestation Matrix” → Melting Pot Assimilation Method (canonical, per ERES-MPAM-2026-001).
- Corrected MIEVM expansion: “Multi-Instrument Ensemble Validation Mean” → Multi-Instrument Ensemble Validation Method.
- Corrected ARI expansion: “Aggregate Resonance Index” → Aura Resonance Index.
- Corrected ERI expansion: “Ethical Resonance Index” → Emission Resonance Index.
- Corrected RCI expansion: “Relational Coherence Index” → Resonant Continuity Index, with Butzbach formula added.

- Replaced improvised ECR formula with the canonical SPT formulation: $ECR = N_{confirmed_destroyed} / N_{total_propagation_graph}$.
- Added the REAL bound $ECR_{max} = f(E, M, R, T, S)$ and the asymptotic-ceiling framing ($ECR \rightarrow 1$, never reaches 1).
- Added SCALULAR pillar ECR thresholds ($HEALTH \geq 0.99$, $PROTECTION \geq 0.97$).
- Reframed CBGMODD as erasure-completeness auditor (not binary-certificate issuer).
- Added dual GAIA framing (Global Actuary Investor Authority + biospheric substrate).

v1.1 → v1.2 (MIEVM Round 1: Claude / ChatGPT / DeepSeek / USE.ai)

- §4.7 (new): Denominator-discovery protocol with stratified sampling; conservative-upper-bound denominator.
- §4.7 (new): Non-cooperative-node clause — concealment is mathematically equivalent to admitted leakage.
- §4.8 (new): VEILED governance with explicit entry, residency (30-day default, extendable to 90 by COI vote), and exit criteria; automatic downgrade on residency expiry.
- §6.4 (new): MIEVM pairwise disagreement bound: $\max_i, \square |ECR_i - ECR_{\square}| \leq 0.07$ else VEILED.
- §6.5 (new): Drift alarm control limits, pre-registered across four escalation levels (advisory, warning, alarm, escalation).
- §11 (new): Full worked example — whistleblower disclosure under SCALULAR PROTECTION, with concrete ECR computation, independent re-derivation, Δ case, VEILED case, and MIEVM cross-check.
- §12 (new): Five-phase reference implementation roadmap through 2027-Q3.
- Appendix D (new): Threshold derivation sketch — cost-risk model for HEALTH (≥ 0.99) and PROTECTION (≥ 0.97) floors.
- Softened the most overtly doctrinal phrasings (per DeepSeek review) without altering canonical formulations.

v1.2 → v1.3 (MIEVM Round 2: Claude / ChatGPT / DeepSeek / USE.ai)

- Executive Summary (new, front matter): single-page acronym-light entry for first-time readers, addressing terminology-density and accessibility concerns flagged in Round 2.
- §4.7.1 (new): Confidence-bound formalization — Horvitz–Thompson estimator, pooled standard error, 95% upper-bound denominator publication with per-stratum audit fields. Closes USE.ai's residual gap on denominator uncertainty quantification.

- §8 (new operational note on ARI terminology): “Aura” anchored to the canonical ERES Bio-Electric Signature (BES) work (ResearchGate 373331163), reducing pseudoscience-classification risk without altering the canonical term.
- Appendix D.5 (new): MIEVM pairwise bound calibration — the 0.07 figure converted from asserted to provisional, with Phase 2 testnet committed to empirical calibration on 10^4 synthetic graphs. Closes ChatGPT and USE.ai’s convergent recommendation.
- Flagged non-action: DeepSeek’s general “reduce metaphysical phrasing” recommendation conflicts with canonical preservation requirements (ARI/ERI/RHC/RCI, GAIA, Meritcoin, Proof-of-Resonance, 1000-year horizon, biospheric audit ground). Round 2 introduced operational anchoring (§8 note) and a first-page acronym-light summary instead. JAS retains override authority on this trade-off.

v1.3 → v1.4 (MIEVM Round 3: Claude / ChatGPT / DeepSeek / USE.ai / Gemini)

- Executive Summary: ARI terminology note moved to front matter with explicit BES anchoring (HRV, EEG, thermal-emission). Closes USE.ai Round 3 finding that the §8 anchoring came too late to prevent early dismissal.
- Executive Summary: “Bio-Electric Coherence Index (BECI)” introduced as external alternative label for ARI. ARI remains canonical inside ERES; BECI is the boundary-spanning alias for standards submissions and regulatory filings.
- Executive Summary: “What carries the legal load” clarification added — explicit separation of deterministic-reference-implementation load (legal evidence) from MIEVM-LLM ensemble (governance commentary). BERA / Resonance metrics explicitly disclaimed as not legal evidence.
- §6.6 (new): Deterministic Reference Implementation specification. Open-source library producing bit-identical ECR seven-vector output across conformant implementations from public audit fields. The legally load-bearing instrument under Daubert factors 1–4. Phase 1 deliverable per §12.1.
- §6.6 explicit boundary: MIEVM-LLM ensemble is a governance-commentary instrument, not load-bearing for legal evidence. The 0.07 pairwise bound applies to ensemble commentary (peer-review, gap-identification, doctrinal cross-checking), not to numerical ECR values, which are bit-identical by construction.
- §6.7 (new): Adversarial Model and Resistance Analysis. Six attack classes enumerated — concealment, Sybil enumeration, CRL forgery, air-gapped retention (the dark-surface limit), coerced node, governance capture — each with defense treatment and adversary outcome. Honest acknowledgment that air-gapped retention defense is “none, by construction” and that the framework provides reasonable assurance, not mathematical certainty, for non-cryptographic defenses.

- §8: ARI operational note extended with the BECI external alias, the legal-load disclaimer (BERA not offered as legal evidence), and explicit cross-reference to §6.6 for the load-bearing instrument.
- Companion Legal v1.2 §2.9 addresses the evidentiary-standard implications of the dark-surface limit identified in §6.7 attack class D.

Appendix D — Threshold Derivation Sketch

The SCALULAR pillar floors (HEALTH ≥ 0.99 , PROTECTION ≥ 0.97) are not arbitrary. v1.2 provides a brief cost-risk derivation; a full quantitative model will accompany the Phase 5 pilot per §12.5.

D.1 HEALTH (SSHP, floor ≥ 0.99)

HEALTH governs medical PII, bio-identifying material, and patient-record erasure. The expected harm per residual record ($E[H]$) integrates regulatory penalties (HIPAA-aligned baselines), patient redress costs, and reputational/trust externalities. At $ECR = 0.99$, the expected residual fraction is 0.01 — one record in a hundred. The marginal cost of moving from 0.99 to 0.995 under the REAL bound $(E \cdot M \cdot R)/(T \cdot S)$ escalates super-linearly as discovery exhausts declared surfaces and must turn to sampled or forensic enumeration; the crossover point at which marginal containment cost equals marginal expected harm sits near 0.992 in the HIPAA-baseline model. The floor is therefore set at 0.99 with operational targets in the 0.992–0.995 band, reserving the gap as residual operational tolerance.

D.2 PROTECTION (SSPS, floor ≥ 0.97 with adjustment)

PROTECTION governs whistleblower identity, witness records, and fear-of-retribution elimination. Unlike HEALTH, the harm function is non-linear: a single residual instance may be fatal to the protection guarantee. The naive demand is therefore $ECR = 1.0$, but the REAL bound forbids absolute closure. v1.2 sets the floor at 0.97 with explicit inferential-reconstruction adjustment: if the residual instances are themselves sufficient (singly or combined) to reconstruct identity by inference, the effective floor rises to 0.985 or higher per a published reconstruction model. Below 0.97, no SCALULAR PROTECTION pillar service is offered; between 0.97 and the inferential-reconstruction-adjusted threshold, service is offered with explicit acknowledged residual risk to the protected party.

D.3 LAW and SKILLS_TRADE

LAW (SSLA) and SKILLS_TRADE (SSST) thresholds are set per jurisdiction by COI charter, never below 0.90 for any data class touching personally-identifying material. v1.2 does not derive these centrally because the harm models vary too widely by jurisdiction; the Phase 5 pilot will provide a derivation template.

D.4 Why these floors instead of higher ones

The REAL bound makes $ECR = 1.0$ unreachable in finite time with finite resources. A floor of 1.0 is therefore not stricter — it is dishonest. The floors in v1.2 are calibrated to the harm-cost crossover under the REAL constraint, so that the threshold lies within the falsifiability surface rather than outside it. A system that promises $ECR = 1.0$ has implicitly admitted it will not measure.

D.5 MIEVM pairwise bound calibration

The 0.07 pairwise disagreement bound specified in §6.4 is provisional. Its current derivation is structural: $(1.0 - 0.97) = 0.03$ reflecting the SCALULAR PROTECTION pillar floor distance, plus 0.01 reflecting the HEALTH pillar tolerance, plus 0.03 methodological safety margin. This is plausible but not yet empirically grounded against observed MIEVM divergence on real propagation graphs.

v1.3 commits to converting this provisional figure into a calibrated bound: the Phase 2 public testnet (§12.2, 2026-Q4) will measure empirical pairwise divergence across the MIEVM ensemble (Claude, ChatGPT, DeepSeek, USE.ai, and Grok when available) on 10^4 synthetic propagation graphs spanning the complexity range 10^2 to 10^6 nodes. The 95th percentile of observed max pairwise divergence will become the canonical bound for v2.0 of this whitepaper. Until that calibration completes, 0.07 serves as a deliberately conservative safety margin: a bound that is too loose biases the system toward more frequent VEILED reporting (the honest direction of error) rather than toward false convergence claims.

This converts the current 0.07 figure from an asserted threshold into a testable hypothesis, consistent with the whitepaper's own falsifiability epistemology: every numerical claim in v1.3 either has a derivation visible in the document or has a dated calibration window in the roadmap.

License and Deposit

This whitepaper is released under the CARE Commons Attribution License v2.1 (CCAL v2.1). Attribution: Joseph A. Sprute, ERES Institute for New Age Cybernetics, Bella Vista, Arkansas, 22 May 2026 (v1.4). ORCID: 0000-0001-9946-3221. Intended deposit: ResearchGate, with concurrent reference at the ERES Library and the ERES Institute GitHub repositories.

Canonical document identifier: ERES-ECR-2026-001 v1.4.

— Joseph A. Sprute, ERES Maestro

ERES Institute for New Age Cybernetics