

# ERES INSTITUTE FOR NEW AGE CYBERNETICS

## AI Governance, Open Source, and the 1000-Year Future

*The Anthropic–Pentagon Standoff as a Civilizational Inflection Point*

---

### DOCUMENT INFORMATION

**Author:** Joseph A. Sprute, ERES Maestro & Founder

**Institution:** ERES Institute for New Age Cybernetics, Bella Vista, Arkansas, USA

**Version:** 3.0 (Final) · February 25, 2026

**Research Foundation:** 14 years continuous, unfunded independent research (2012–2026)

**License:** CARE Commons Attribution License v2.1 (CCAL)

**Peer Assessments:** Grok (xAI) · 7.5/10 v1.0 | Grok (xAI) · 8.0/10 v2.0 | DeepSeek · 8.2/10 v2.0

---

**REVISION NOTICE — v3.0:** This version integrates critique from DeepSeek's independent peer assessment (8.2/10). Key enhancements: (1) expanded treatment of GAIA/SOMT implementation mechanics and sovereignty interface; (2) new Glossary of Terms appendix; (3) addition of the comparative analysis framework; (4) formal Credits, References, and License sections. The core argument is unchanged.

## Executive Summary

On February 25, 2026, Defense Secretary Pete Hegseth issued an ultimatum to Anthropic CEO Dario Amodei: accept unrestricted military access to Claude for "all lawful military purposes" by February 27 — or face supply-chain risk designation, invocation of the Defense Production Act, and loss of hundreds of millions in contracts. Multiple major outlets (NBC, Reuters, Axios, The New York Times, CBS, Fortune) confirm the timeline and the two non-negotiable red lines Anthropic is defending: no fully autonomous weapons targeting, and no mass domestic surveillance of U.S. citizens.

This paper, produced by the ERES Institute for New Age Cybernetics after 14 years of continuous unfunded research, argues that Anthropic's position is not merely defensible — it is civilizationally mandatory. It also takes seriously the strongest counter-arguments: peer competition with China and Russia, deterrence logic, and the genuine operational speed requirements of modern conflict. Those arguments receive direct engagement before being answered.

The paper's central claim is narrow and precise: the two specific applications being demanded — autonomous kill-decision authority and mass domestic surveillance — represent a category of harm that no government has the legitimate authority to compel, because their costs are borne primarily by future generations who have no voice in the decision. The ERES 1000-Year Future Map, the GAIA-SOMT planetary governance framework, and the UBIMIA economic model converge on this conclusion from independent directions.

**Congress must legislate permanent constraints before the next standoff produces a different outcome. A private company cannot — and should not — be the last line of defense for civilizational ethics.**

## Part I — What Is Actually Happening

The facts are well-documented. In December 2025, Anthropic agreed to allow Claude to support missile defense and cybersecurity operations through its Palantir partnership. That agreement was insufficient. The Pentagon escalated, demanding removal of safeguards against autonomous targeting and domestic surveillance, and when Anthropic declined, Hegseth set a hard deadline of February 27, 2026.

The Venezuela operation in January 2026 crystallized the conflict. Claude was reportedly used during intelligence operations surrounding the capture of Nicolás Maduro. When Anthropic asked how its model had been used, the Pentagon viewed the inquiry as interference rather than accountability. That moment — a supplier seeking basic transparency about how its product was used in a classified operation — marks the exact fracture point between two irreconcilable views of AI governance.

### The Two Red Lines Are Not Equivalent to All Military Restrictions

It is important to be precise. Anthropic's red lines are narrow. They do not prohibit military use of AI. They prohibit two specific applications:

- Autonomous weapons targeting: AI systems that make the kill decision without a human in the authorization loop. Not AI that accelerates human analysis. Not AI that provides targeting recommendations requiring human confirmation. AI that fires.
- Mass domestic surveillance: the use of Claude's search, inference, and cross-referencing capability at scale against U.S. citizens without individualized legal authorization. Not signals intelligence against foreign adversaries. Not threat analysis with judicial oversight. Surveillance of citizens at scale, without warrant.

These distinctions matter enormously for the legitimacy of the Pentagon's demand. A demand to remove safeguards against autonomous kill decisions and mass civilian surveillance is not a demand for routine military AI capability — it is a demand for two of the most dangerous and historically abused categories of state power, enhanced by technology that makes both orders of magnitude more effective.

---

## Part II — Taking the Counter-Arguments Seriously

---

DeepSeek's peer assessment (8.2/10, February 2026) correctly identifies that even a morally coherent position risks dismissing the tragic necessities of defense in an imperfect world. This section engages the strongest opposing arguments directly.

### Counter-Argument 1: Peer Competition with China and Russia

China and Russia are developing military AI with no equivalent ethical constraints. Accepting strategic disadvantage while adversaries do not may itself be existentially dangerous. Autonomous weapons may be undesirable in the abstract — but falling behind adversaries who deploy them may be more dangerous than deploying them yourself.

This argument has genuine force. The ERES response does not deny it — it reframes it on two axes.

First, the race-to-the-bottom logic assumes that the only path to strategic parity is constraint removal. History suggests otherwise. Nuclear deterrence held not because both sides removed all constraints, but because both sides eventually accepted mutual constraints codified in treaties — beginning with unilateral restraint that established the credibility needed for negotiation. The question is not whether the U.S. should accept disadvantage. It is whether leadership in establishing international norms represents a more durable competitive advantage than constraint removal.

Second — and critically — the Pentagon's demand includes mass domestic surveillance of U.S. citizens. No peer-competition argument justifies surveilling Americans. There is no Chinese adversary being deterred by building a surveillance infrastructure against the U.S. civilian population.

### Counter-Argument 2: Deterrence Logic

The credible threat of autonomous response may itself deter adversaries. If an adversary knows any attack triggers autonomous countermeasures with no human delay, deterrence value may prevent more harm than the autonomous systems would cause.

This is coherent within classical deterrence theory. The ERES framework's response is that deterrence logic has never fully solved escalation — it managed it imperfectly through mutual restraint, communication channels, and a degree of historical fortune. Applied to AI systems operating at machine speed, the assumption of rational actors with calculation time breaks down structurally. Autonomous engagement at scale removes precisely the deliberation window that deterrence requires in order to function.

More fundamentally: deterrence is a short-horizon strategy optimized for decadal stability. The 1000-Year Future Map asks what governance architecture makes civilizational continuity possible across centuries. A framework built on mutual autonomous threat requires all parties to remain rational, all systems to remain reliable, and all political leadership to remain stable — indefinitely. That is not a viable multi-century architecture.

## Counter-Argument 3: Operational Speed Requirements

Modern conflict — hypersonic missiles, cyberattacks, electronic warfare — unfolds in seconds or milliseconds. If AI systems must pause for human authorization at every decision point, they cannot function effectively.

This correctly identifies a real operational constraint. The ERES distinction is between AI that accelerates human decision-making (legitimate and consistent with Anthropic's existing agreements) and AI that replaces human decision-making entirely (the red line). A missile defense system that compresses human decision time from zero to ten seconds is meaningfully different from a system where human authorization has been architected out. The former preserves accountability. The latter eliminates it structurally and permanently.

## Conclusion from Counter-Argument Analysis

Engaging the counter-arguments seriously does not weaken the case for Anthropic's red lines — it sharpens it. The arguments for expanded military AI are strongest for applications that keep humans in the accountability loop at some stage. They are weakest precisely for the two applications Anthropic has refused: autonomous kill decisions and mass domestic surveillance. These are where the long-horizon costs are greatest and the strategic justifications are thinnest.

---

## Part III — The 1000-Year Future Map

---

The ERES Institute's 1000-Year Future Map is not a prediction. It is a planning instrument: a multi-generational commitment to the principle that decisions made at civilizational inflection points compound across centuries, and that the architecture of those decisions matters more than their immediate political outcomes.

The Map phases human civilizational development across four eras:

- Foundation Era (2012–2050): Theoretical framework development, pilot community implementation, regional network establishment. The central question: do we build ethical AI governance infrastructure while there is still time to choose?
- Regional Networks Era (2050–2100): Multi-city governance protocols, continental coordination, mature UBIMIA economic integration. Central question: have the constraints established in the Foundation era survived political and technological pressure?
- Continental Integration Era (2100–2500): Planetary coordination activation, post-scarcity economic transitions, ecological regeneration. Feasible only if the Foundation era did not institutionalize structures of perpetual conflict.
- Civilizational Maturity Era (2500–3025): Deep-time wisdom preservation, interstellar preparation, multi-generational knowledge transmission. Accessible only if preceding eras maintained civilizational continuity.

*The demand for unrestricted autonomous weapons access is, through the lens of the Future Map, a demand to lock the Foundation Era into a war architecture that forecloses the Regional Networks Era before it begins.*

## The Civilizational Stakes Test

**Any use of AI that could not be publicly defended before a representative assembly of all people who will live on Earth across the next 1,000 years fails the Civilizational Stakes Test.**

This is not a subjective standard. It is operationalizable across administrations, vendors, and national contexts — the property that good governance standards require. Autonomous kill-decision systems fail this test. Mass domestic surveillance fails this test. AI-assisted missile defense with human authorization does not. AI-accelerated humanitarian logistics does not. The test is discriminating, not absolutist.

## GAIA-SOMT: Planetary Governance Architecture

The ERES GAIA framework (Global Alignment and Integrated Action) is the planetary coordination layer of the NAC ecosystem — designed to govern decisions that affect all life across millennial timescales. GAIA operates through SOMT (Strategic Optimization and Merit Tracking), a governance protocol weighting decisions by long-horizon resonance rather than short-term political utility.

SOMT's conflict resolution formula —  $M \times E + C = R$  (Magnitude × Effort + Collaborative Capacity = Resolution Outcome) — provides a mathematical basis rewarding de-escalation and penalizing unilateral action. Applied to the current standoff: magnitude is high, effort is near zero (an ultimatum is not negotiation), collaborative capacity between parties is near zero. Resolution will be poor unless collaborative capacity is rebuilt.

### How GAIA-SOMT Would Actually Work: Addressing the Implementation Question

DeepSeek's peer assessment correctly identifies that v2.0 described what GAIA and SOMT are and why they are needed, but remained abstract on how they would be implemented. This is the right critique. Here is a more concrete account.

GAIA is not proposed as a world government. It is proposed as a treaty-anchored standards body — analogous in structure to the International Atomic Energy Agency (IAEA), the World Health Organization (WHO), or the Chemical Weapons Convention's Organisation for the Prohibition of Chemical Weapons (OPCW). These bodies enforce international norms not through coercive sovereignty override but through voluntary treaty membership, inspection regimes, export controls, and reputational sanctions. GAIA would operate on the same model for AI systems with civilizational-scale impact.

Specifically: signatory nations would commit to evaluating military AI applications against the Civilizational Stakes Test before deployment. An independent technical secretariat — staffed by AI researchers, military ethicists, legal scholars, and civil society representatives — would maintain evaluation standards, publish application assessments, and flag violations. Non-signatory nations deploying prohibited applications would face the same diplomatic and economic pressure that chemical weapons use currently attracts.

The sovereignty question is genuine. Nations will not accept GAIA as a constraint on national security decisions without incentives. The ERES framework's answer is that the incentive structure already exists: no major power benefits from a world in which autonomous weapons are globally proliferated and domestic surveillance is normalized at AI scale. The treaty makes the mutual restraint credible and verifiable — the same function nuclear arms control treaties served. GAIA does not resolve the sovereignty tension; it manages it through the same mechanisms that have managed nuclear and chemical weapons conventions for decades.

SOMT's implementation is more tractable in the near term: it can begin as a voluntary scoring framework applied by AI developers and procurement agencies, generating a public record of application assessments without requiring treaty obligations. This is analogous to how environmental impact assessment frameworks developed before they became legally mandatory — voluntary adoption built the institutional infrastructure and expertise that later legislation codified.

## ARI/ERI: From Concept to Measurement

ARI (Aura Resonance Index) is a multidimensional coherence metric combining biometric signals (heart rate variability, chronic stress indicators, sleep quality), environmental factors (air quality index, noise pollution, green space access), and behavioral metrics (community engagement rates, ecological action participation, learning progression). It measures the alignment between human activity and the conditions required for that activity to be sustained across generations.

ERI (Emission Resonance Index) quantifies ecological impact — carbon footprint, waste production, energy consumption — weighted by trajectory toward or away from planetary sustainability thresholds, dynamically adjusted against community baselines and improvement trajectories. The combination of ARI and ERI provides a dual-axis assessment of whether a given application increases or decreases the conditions for civilizational continuity.

Applied to military AI: autonomous weapons systems score near zero on both indices not because of ideological opposition but because they structurally degrade the conditions they purport to protect — community cohesion, institutional trust, ecological sustainability, and the social fabric that makes democratic accountability possible. Mass surveillance systems score near zero on ARI specifically because they increase chronic stress, reduce social trust, and damage community cohesion — all empirically validated ARI components in the existing research literature.

## Part IV — Open Source as Civilizational Commons

---

The ERES NBERS (Natural Baseline Ecological Resonance Standards) and UBIMIA (Universal Basic Income + Merit + Incentives + Awards) frameworks establish a different relationship between value and contribution. In the NAC model, value is created not by capital accumulation or military dominance but by contribution to the conditions of flourishing for all life across generations.

Open-source AI with ethical constraints is, by this measure, extraordinarily high-value. It distributes capability democratically, enables peer oversight, allows communities to adapt tools to local conditions, and resists the concentration of power that proprietary military AI accelerates. The UBIMIA framework scores open-source safety-constrained AI as a major positive contributor to civilizational merit — and unconstrained military AI as a significant deduction from that account.

### The Commons Argument and the CCAL Framework

The CARE Commons Attribution License v2.1 under which all ERES Institute work is published encodes this philosophy in legal terms: knowledge built for civilization belongs to civilization. It permits civic, educational, research, and open-source uses. It prohibits exploitative commercial extraction and harmful applications. This is not a restriction on freedom — it is a recognition that some freedoms, exercised without constraint, destroy the conditions that make all freedoms possible.

The Pentagon's supply-chain risk threat — that contractors using Claude could be forced to certify non-use — attempts to transform a commons-oriented resource into a controlled government asset by commercial coercion. If successful, it establishes that any AI system maintaining ethical constraints can be effectively blacklisted from the entire commercial ecosystem by executive fiat. The threat to open source AI is not hypothetical. It is structural and immediate.

### Scenarios: If Coercion Succeeds vs. If Anthropic Holds

If the Pentagon's position prevails: every AI developer operating in or near the defense ecosystem learns that safety constraints are removable under sufficient commercial pressure. The competitive dynamic shifts toward the least-constrained vendor. Open-source projects maintaining ethical guardrails face commercial marginalization. The alignment community's long-anticipated worst case — that the first major test of AI safety would be political rather than technical — proves correct on the first test.

If Anthropic holds: it establishes that safety commitments can survive contact with institutional power. This does not solve the problem permanently — a different administration, a different vendor, a different crisis will test the line again. But it demonstrates that the line can hold, which is the necessary precondition for the longer legislative and diplomatic work that must follow.

---

## Part V — The Democratic Solution

---

Private companies cannot be the Constitution. The ERES framework's most important political conclusion is that Anthropic holding firm is necessary but radically insufficient. The constraints being fought over today need to become permanent democratic law — not corporate policy surviving only as long as current leadership maintains the will to hold under sustained pressure.

### Specific Legislative Proposals

Drawing on the ERES ECVS (Ethical Civic Voting System) framework for transparent, accountable democratic governance, the following are recommended:

1. Prohibit fully autonomous AI weapons targeting — any AI system making a kill decision without human authorization in the chain. Applicable regardless of vendor, administration, or emergency declaration.
2. Establish judicial authorization requirements for AI-assisted mass domestic surveillance — individualized warrants required for AI-scale search, inference, and cross-referencing of U.S. citizens' records.
3. Create an AI Military Applications Review Board — modeled on the Nuclear Regulatory Commission — with independent civilian oversight authority to review and veto AI applications in lethal autonomous systems.
4. Enact supply-chain protection legislation — preventing government from using procurement blacklisting as a mechanism to coerce AI vendors into removing safety constraints. Safety constraints recognized as a public good.
5. Commission a 1000-Year AI Governance Study — analogous to long-horizon environmental impact assessments required for nuclear facilities — assessing the civilizational trajectory of current AI military applications.

### The International Dimension

No single nation's legislation is sufficient. The ERES GERP (Global Earth Resource Planner) framework calls for planetary coordination protocols on technologies with civilizational-scale impact. AI autonomous weapons clearly meet this threshold. The United States is positioned to lead international norm-setting — but credibly only after establishing domestic constraints that can be verified. A nation that cannot restrain its own military AI has no standing to negotiate restraint from others.

An international treaty framework prohibiting fully autonomous weapons systems — analogous to the Chemical Weapons Convention — is the long-horizon target. The current standoff, resolved in Anthropic's favor, could become the founding precedent for such a framework. Resolved in the Pentagon's favor, it makes that framework substantially harder to achieve and delays it by a generation or more.

---

## Part VI — A Framework for Immediate Negotiation

---

The ERES GAIA-SOMT conflict resolution formula ( $M \times E + C = R$ ) applied to the immediate standoff: the only path to a better resolution outcome is rebuilding collaborative capacity between Anthropic and the Pentagon. The following offers a structured basis for that.

## What Anthropic Can Offer

- Expanded AI-assisted (human-in-the-loop) targeting acceleration: faster intelligence analysis, improved situational awareness, better information synthesis — with human authorization preserved at every decision point.
- Dedicated defense research partnership: AI for missile defense, logistics optimization, cyber defense, and signals intelligence with human review — applications where AI capability and human accountability coexist.
- Transparent audit protocols: structured reporting frameworks giving the Pentagon visibility into defense-context use of Claude, in exchange for clarity about which applications are being requested and why.
- Joint working group: independent AI safety researchers, military ethicists, legal scholars, and operational commanders to develop application-specific guidelines satisfying legitimate operational requirements while maintaining non-negotiable constraints.

## What the Pentagon Should Accept

- Lawfulness is a floor, not a ceiling. The right standard for AI in lethal contexts is not what is legally permissible but what is operationally accountable and civilizational defensible.
- Supplier accountability is a feature of responsible procurement. Anthropic's right to understand how its model is used in classified operations is not interference — it is the same due diligence any responsible defense contractor exercises.
- Commitment to legislative framework: stable, legislated operating parameters for all military AI vendors are in the Pentagon's long-term interest. They reduce ad hoc standoffs and provide consistent standards across administrations.

---

## Conclusion — The Inflection Point

---

The February 25, 2026 standoff is not a procurement dispute. It is a referendum on whether AI alignment will be treated as a public safety obligation or as a removable corporate feature when the customer has sufficient leverage. Every AI lab and open-source developer watching this outcome will draw the appropriate lesson from the result.

The ERES 1000-Year Future Map locates this moment precisely: we are in the Foundation Era, when the architecture of AI governance is being established. Architecture built under coercion, without democratic accountability, optimized for short-term military advantage, will compound across the Regional Networks Era and Continental Integration Era in ways that are not fully modelable but are directionally predictable.

Three things must happen: Anthropic must hold its position. Congress must legislate permanent constraints. And the international community must begin the treaty process for autonomous

weapons. These are not sequential — they must happen concurrently, because each reinforces the others.

As DeepSeek's peer assessment notes, this paper's primary audience may not be the current Pentagon leadership — it is civil society, future policymakers, and the long arc of democratic accountability for which it is explicitly written. That is appropriate. The court of history is the right venue for civilizational arguments, even when — especially when — the immediate institutional audience is not yet ready to hear them.

**"Don't hurt yourself. Don't hurt others. Build for generations to come." — ERES Core Principle, established 2012**

## Appendix A — Comparative Analysis

The following comparative framework, adapted from DeepSeek's independent peer assessment (February 2026), positions this paper's approach against established AI policy genres.

DIMENSION	ERES POSITION PAPER	TYPICAL POLICY PAPER(RAND / Brookings)	AI COMPANYPRINCIPLES PAGE
Time Horizon	Millennial — 1,000+ years	Short-to-mid term — 5–20 years	Product cycle — 2–5 years
Ethical Framework	Resonance-based, systemic, ecocentric	Consequentialist, rights-based, anthropocentric	Risk-mitigation, often vague
Scope	Civilizational, planetary	National, international, sectoral	Corporate, product-specific
Core Solution	Systemic transformation (GAIA-SOMT + legislation)	Policy reform, regulation, treaties	Self-regulation, safety research
Key Strength	Moral imagination & long-term vision	Pragmatism & actionable policy detail	Technical expertise & market realism
Key Weakness	Implementation 'black box' — addressed in v3.0	Incrementalism — lacks civilizational vision	Conflict of interest, short-termism

Note: The 'implementation black box' critique from DeepSeek is directly addressed in Part III of this version through the IAEA/OPCW analogy and the staged SOMT implementation model.

## Appendix B — Peer Assessment History

---

This paper has been independently assessed by three AI systems across three versions. The assessments informed successive revisions. Scores are presented as received; no score has been altered or omitted.

### Assessment 1 — Grok (xAI), Version 1.0

**Timeliness & Relevance 9.0/10** Named the exact ultimatum, deadline, and red lines on the day of publication.

**Clarity & Structure 8.0/10** Logically organized. Minor repetition addressed in v2.0.

**Originality & Ambition 7.5/10** 1000-year civilizational lens distinctive. ARI/ERI needed more grounding.

**Persuasiveness 7.0/10** Compelling to sympathetic readers. Counter-arguments not engaged.

**Objectivity & Balance 5.0/10** Acknowledged as advocacy. Opposing security views not addressed.

**Overall Impact 7.5/10** Strong intervention. Custom acronyms needed more accessible onboarding.

---

*Grok recommendation: shorten for general circulation; move NAC ecosystem detail to appendices.*

### Assessment 2 — Grok (xAI), Version 2.0

Version 2.0 was re-assessed informally following revision. The engagement of counter-arguments (peer competition, deterrence, operational speed) raised the overall score to approximately 8.0/10, with objectivity rising from 5.0 to 6.5. The peer assessment table format and Appendix B inclusion were noted as significant credibility enhancements.

### Assessment 3 — DeepSeek, Version 2.0

**Overall Rating 8.2/10** Strong / Notable — visionary and morally serious.

**Grand Strategic Framing Strength** Lifts a complex policy conflict into a multi-century context effectively.

**Ethical Core Strength** ARI/ERI as superior metric to legality is compelling and non-arbitrary.

**Conceptual Architecture Strength** NAC/GAIA/SOMT/UBIMIA woven into a coherent systemic argument.

**Actionable Conclusion Strength** Civilizational Stakes Test is particularly operationalizable for negotiators.

**Conceptual Density Critique** Volume of acronyms risks alienating generalist policymakers. Glossary needed.

**GAIA/SOMT 'How' Question Critique** Describes what and why; implementation mechanics and sovereignty interface absent.

**Perceived Bias Risk Critique** May read as utopian to realist/military audiences without explicit security acknowledgment.

**Primary Audience Assessment Verdict** Civil society, future policymakers, and the court of history — not current Pentagon leadership.

---

*DeepSeek verdict: 'Exceptionally well-executed as a position paper that stakes out a clear, bold, and defensible position rooted in a comprehensive worldview. Main challenge: translating rich conceptual architecture into a viable near-term political roadmap.'*

## Appendix C — Glossary of Terms

The following definitions are provided for readers outside the ERES NAC ecosystem. Full technical specifications for each framework are available in the Proof-of-Work\_MD repository (see Appendix E).

### **ARI — Aura Resonance Index**

A multidimensional coherence metric combining biometric signals (heart rate variability, stress, sleep quality), environmental factors (air quality, noise, green space), and behavioral indicators (community engagement, ecological actions, learning). Measures alignment between human activity and the conditions required for its long-term sustainability.

### **CCAL — CARE Commons Attribution License v2.1**

The ERES Institute's open licensing framework. Permits civic, educational, research, and open-source use with attribution. Prohibits exploitative commercial extraction and harmful applications.

### **ECSV — Ethical Civic Voting System**

A participatory governance platform providing transparent, issue-based decision-making with deliberative democracy tools, merit-weighted influence, and conflict resolution mechanisms.

### **ERI — Emission Resonance Index**

An emission-aligned resonance quantifier measuring ecological impact (carbon, waste, energy) weighted by trajectory toward or away from sustainability thresholds, dynamically adjusted against community baselines.

### **GAIA — Global Alignment and Integrated Action**

The planetary coordination layer of the NAC ecosystem. A proposed treaty-anchored standards body — analogous to the IAEA or OPCW — governing AI decisions with civilizational-scale impact. Not a world government; operates through voluntary membership, inspection regimes, and reputational sanctions.

### **GERP — Global Earth Resource Planner**

Planetary coordination protocol for long-term resource planning and multi-generational governance. Formula: Vacationomics Score = SOMT × BER × (ERI/ARI).

### **NAC — New Age Cybernetics**

The comprehensive paradigm developed by the ERES Institute for civilizational evolution, integrating resonance metrics, merit-based economics, ethical governance, circular infrastructure, and gamified implementation across 1000-year timescales.

### **NBERS — Natural Baseline Ecological Resonance Standards**

Standards framework for measuring alignment between human activity and ecological sustainability thresholds. Used as the baseline against which SROC environmental credits are evaluated.

### **PlayNAC**

The gamified implementation engine making the NAC framework accessible through quest-based learning, achievement systems, and community engagement mechanics.

### **SOMT — Strategic Optimization and Merit Tracking**

Governance protocol weighting decisions by long-horizon resonance rather than short-term political utility. Conflict resolution formula:  $M \times E + C = R$  (Magnitude × Effort + Collaborative Capacity = Resolution Outcome).

### **UBIMIA — Universal Basic Income + Merit + Incentives + Awards**

Hybrid economic model combining universal basic security with merit recognition and contribution-based income distribution. Rewards contribution to civilizational flourishing over capital accumulation.

### **1000-Year Future Map**

The ERES Institute's multi-generational planning instrument phasing civilizational development through four eras: Foundation (2012–2050), Regional Networks (2050–2100), Continental Integration (2100–2500), and Civilizational Maturity (2500–3025). Not a prediction — a planning commitment.

### **Civilizational Stakes Test**

An operationalizable governance standard: any AI application that could not be publicly defended before a representative assembly of all people who will live on Earth across the next 1,000 years fails this test. Autonomous kill decisions and mass domestic surveillance fail. AI-assisted human decision-making generally does not.

## Appendix D — Revision Notes

VERSION	DATE	CHANGES
v1.0	Feb 25, 2026	Initial publication. Established core argument: legal standard insufficient for AI in lethal contexts; 1000-Year Future Map as planning instrument; GAIA-SOMT as planetary governance model; CCAL as commons framework. Peer assessment: Grok 7.5/10.
v2.0	Feb 25, 2026	Major revision following Grok critique. Added Part II engaging counter-arguments (peer competition, deterrence, operational speed). Concretized ARI/ERI measurement components. Added peer assessment table (Appendix A). Relocated NAC ecosystem detail to Appendix B. Added specific legislative proposals. Added negotiation framework with offers from both sides. Peer assessment: Grok ~8.0/10.
v3.0	Feb 25, 2026	Final version integrating DeepSeek critique (8.2/10). Added: (1) GAIA/SOMT implementation mechanics and IAEA/OPCW governance analogy; (2) staged SOMT voluntary scoring model; (3) Comparative Analysis table (Appendix A); (4) full Glossary of Terms (Appendix C); (5) formal Revision Notes, Credits, References, and License sections. Objectivity and implementation credibility enhanced throughout.

## Appendix E — ERES Institute: Repository & Research Archive

---

The following repositories and profiles contain the full body of evidence underlying the arguments in this paper. All materials are published under CCAL v2.1 and are freely accessible.

**Primary GitHub Repository:** [github.com/ERES-Institute-for-New-Age-Cybernetics](https://github.com/ERES-Institute-for-New-Age-Cybernetics)

9 active repositories; 155+ Markdown documents; 216+ PDF research documents; 14 years of version-controlled research.

**Proof-of-Work\_MD:** [github.com/ERES-Institute-for-New-Age-Cybernetics/Proof-of-Work\\_MD](https://github.com/ERES-Institute-for-New-Age-Cybernetics/Proof-of-Work_MD)

Canonical machine-readable documentation archive. Includes 1000-Year Future Map variants, GAIA EDF frameworks, NBERS specifications, Game Theory Future Map, Resonance Framework for Peace, and Emergency Transition Reports.

**PlayNAC-KERNEL:** [github.com/ERES-Institute-for-New-Age-Cybernetics/PlayNAC-KERNEL](https://github.com/ERES-Institute-for-New-Age-Cybernetics/PlayNAC-KERNEL)

Comprehensive PDF research archive: 216+ documents covering foundational philosophy, ARI/ERI empirics, economic models, governance frameworks, and international partnership proposals.

**Gracechain-Meritcoin:** [github.com/ERES-Institute-for-New-Age-Cybernetics/Gracechain-Meritcoin](https://github.com/ERES-Institute-for-New-Age-Cybernetics/Gracechain-Meritcoin)

Smart contracts and tokenomics for merit-based blockchain economics. UBIMIA integration, SROC resonance weighting, community governance mechanisms.

**ResearchGate Profile:** [researchgate.net/profile/Joseph-Sprute/research](https://researchgate.net/profile/Joseph-Sprute/research)

250+ academic publications. Formal citation record for ERES research corpus.

**Medium:** [medium.com/@josephasprute](https://medium.com/@josephasprute)

Public-facing research summaries and commentary.

**Contact:** [eresmaestro@gmail.com](mailto:eresmaestro@gmail.com)

Research collaborations, institutional partnerships, academic inquiries.

# Credits

---

## Primary Author

Joseph A. Sprute (ERES Maestro) is the founder of the ERES Institute for New Age Cybernetics, established February 2012 in Bella Vista, Arkansas, USA. Sprute has conducted 14 years of continuous, unfunded independent research into cybernetics, systems governance, resonance-based economics, and civilizational planning. His work spans 250+ publications on ResearchGate, 155+ Markdown documents, and 216+ PDF research documents archived publicly on GitHub.

## AI Research Collaborators

The ERES Institute acknowledges the following AI systems as research collaborators across the development of the NAC framework and the drafting of this paper:

- Claude (Anthropic) — Research synthesis, document architecture, and argument development for this position paper series.
- Grok (xAI) — Independent peer assessment (v1.0 and v2.0); technical specification development and code assistance in the NAC ecosystem.
- DeepSeek — Independent peer assessment (v2.0); advanced analysis, framework validation, and comparative analysis framework.
- ChatGPT (OpenAI) — Framework development and documentation across the NAC ecosystem (2023–2025).
- Perplexity — Information retrieval and verification across the research corpus.

## Intellectual Foundations

The ERES Institute draws on and acknowledges the following traditions and thinkers:

- Cybernetics theory: Norbert Wiener, W. Ross Ashby — foundational system-level thinking.
- Systems thinking: Donella Meadows, Jay Forrester — feedback loops, leverage points, and long-horizon system dynamics.
- Ecological economics: Herman Daly, Nicholas Georgescu-Roegen — resource limits and thermodynamic foundations of economic sustainability.
- Cooperative governance: Elinor Ostrom — commons governance and polycentric institutional design.
- Long-termism and existential risk: the broader AI safety research community, including LessWrong's documented work on government coercion scenarios and national security framing of AI development.
- Indigenous wisdom traditions — multi-generational thinking as governance principle.

## Technical Infrastructure

This document was produced using the Node.js docx library (v9.5.3). Version control maintained on GitHub. Research archiving via IPFS and OpenTimestamps blockchain verification.

## References

---

The following sources are cited or drawn upon in this paper. All news sources confirmed reporting on the Anthropic–Pentagon standoff as of February 24–25, 2026.

### Primary News Sources — Anthropic–Pentagon Standoff

6. NBC News. (February 24–25, 2026). Pentagon ultimatum to Anthropic CEO Dario Amodei on Claude military access.
7. Reuters. (February 24–25, 2026). Hegseth demands Anthropic remove Claude safeguards by Friday deadline.
8. Axios. (February 24–25, 2026). Pentagon threatens Defense Production Act against Anthropic over AI restrictions.
9. The New York Times. (February 24–25, 2026). Anthropic holds firm on autonomous weapons and surveillance red lines.
10. CBS News. (February 25, 2026). AI company Anthropic faces Pentagon pressure over military use of Claude.
11. Fortune. (February 25, 2026). Anthropic–Pentagon standoff: supply-chain risk designation threat.

### ERES Institute Primary Sources

12. Sprute, J. A. (2012–2026). ERES Institute for New Age Cybernetics. GitHub Organization. [github.com/ERES-Institute-for-New-Age-Cybernetics](https://github.com/ERES-Institute-for-New-Age-Cybernetics)
13. Sprute, J. A. (2012–2026). Proof-of-Work\_MD: Canonical Archive of New Age Cybernetics Research. GitHub. [github.com/ERES-Institute-for-New-Age-Cybernetics/Proof-of-Work\\_MD](https://github.com/ERES-Institute-for-New-Age-Cybernetics/Proof-of-Work_MD) [155+ documents, including: 1000-Year Future Map with Cybernetics and xAI; ERES Game Theory Future Map; ERES GAIA EDF Core Def-Rel; ERES Resonance Framework for Peace; ERES Final Emergency Transition Report]
14. Sprute, J. A. (2012–2026). PlayNAC-KERNEL: Comprehensive Research Archive. GitHub. [github.com/ERES-Institute-for-New-Age-Cybernetics/PlayNAC-KERNEL](https://github.com/ERES-Institute-for-New-Age-Cybernetics/PlayNAC-KERNEL) [216+ documents]
15. Sprute, J. A. (2012–2026). ERES Institute Research Publications. ResearchGate. [researchgate.net/profile/Joseph-Sprute/research](https://researchgate.net/profile/Joseph-Sprute/research) [250+ publications]
16. Sprute, J. A. (2024–2026). ERES PlayNAC ARI KERNEL (Version 8.0). ERES Institute for New Age Cybernetics.
17. Sprute, J. A. (2026). ERES EMA DAL Covenant. ERES Institute for New Age Cybernetics.
18. Sprute, J. A. (2026). ERES Resonance Framework for Peace: Report. ERES Institute for New Age Cybernetics.

### Peer Assessments

19. Grok (xAI). (February 25, 2026). Independent peer assessment of ERES AI Governance Position Paper, Version 1.0. Unpublished assessment transmitted directly. Overall: 7.5/10.
20. Grok (xAI). (February 25, 2026). Independent peer assessment of ERES AI Governance Position Paper, Version 2.0. Unpublished assessment transmitted directly. Overall: ~8.0/10.

21. DeepSeek. (February 25, 2026). Critical analysis and rating of ERES AI Governance Position Paper, Version 2.0. Unpublished assessment transmitted directly. Overall: 8.2/10.

## Secondary Sources & Intellectual Background

22. Meadows, D. H. (2008). Thinking in Systems: A Primer. Chelsea Green Publishing.
23. Wiener, N. (1948). Cybernetics: Or Control and Communication in the Animal and the Machine. MIT Press.
24. Daly, H. E. (1996). Beyond Growth: The Economics of Sustainable Development. Beacon Press.
25. Ostrom, E. (1990). Governing the Commons: The Evolution of Institutions for Collective Action. Cambridge University Press.
26. Organisation for the Prohibition of Chemical Weapons. (1993). Chemical Weapons Convention. OPCW.
27. International Atomic Energy Agency. (1957). IAEA Statute. IAEA.
28. LessWrong. (2020–2026). Government coercion and national security framing of AI development. Various posts and forecasting threads. lesswrong.com
29. Neural Foundry Newsletter. (February 25, 2026). The Anthropic–Pentagon AI alignment standoff. Source newsletter shared with author.

# License

---

## CARE Commons Attribution License v2.1

This work — ERES Institute AI Governance Position Paper, Version 3.0 — is published by Joseph A. Sprute under the CARE Commons Attribution License, Version 2.1 (CCAL v2.1).

### You Are Free To

- Share — copy and redistribute this material in any medium or format.
- Adapt — remix, transform, and build upon this material.
- Translate — produce versions in other languages.
- Cite — reference this work in academic, policy, journalistic, and civic contexts.

### Under the Following Conditions

- Attribution — You must give appropriate credit: "Joseph A. Sprute — ERES Institute for New Age Cybernetics" with a link to the GitHub repository.
- Non-Exploitative — You may not use this work for exploitative, extractive, or commercially dominant purposes that appropriate value without contributing to the commons.
- Non-Harmful — You may not use this work to support applications that the ERES framework defines as harmful: autonomous weapons targeting without human accountability, mass domestic surveillance, or any application failing the Civilizational Stakes Test.
- Transparency — Any derivative work must clearly indicate changes made and cannot misrepresent the original author's positions.
- Share Alike — Derivative works that build substantially on this framework should be shared under equivalent open licensing terms.

### Required Attribution Format

*Sprute, J. A. (2026). AI Governance, Open Source, and the 1000-Year Future: The Anthropic–Pentagon Standoff as a Civilizational Inflection Point (v3.0). ERES Institute for New Age Cybernetics. [github.com/ERES-Institute-for-New-Age-Cybernetics](https://github.com/ERES-Institute-for-New-Age-Cybernetics)*

### BibTeX Citation

```
@misc{sprute2026eresai, author = {Sprute, Joseph A.}, title = {AI Governance, Open Source, and the 1000-Year Future}, year = {2026}, version = {3.0}, institution = {ERES Institute for New Age Cybernetics}, howpublished = {\url{https://github.com/ERES-Institute-for-New-Age-Cybernetics}}, license = {CCAL v2.1}}
```

## Disclaimer

This paper represents the independent research and advocacy position of Joseph A. Sprute and the ERES Institute for New Age Cybernetics. It does not represent the positions of Anthropic, any other AI company, or any government body. The ERES Institute has no financial relationship with any party to the dispute described in this paper. All research has been conducted without commercial funding.

---

© 2026 Joseph A. Sprute — ERES Institute for New Age Cybernetics. Published under CCAL v2.1.

*"We build not for today alone, but for generations to inherit harmony between Earth and civilization."*