ERES INSTITUTE FOR NEW AGE CYBERNETICS
ResearchGate Publication #370

The Transparency Deficit:

Why Closed AI Systems Require Open Reference Architectures

and How ERES Provides the Missing Governance Layer

A Case Study in Structural Fragility, Semantic Governance,
and the Anthropic Claude Code Leak of March 31, 2026

Joseph Allen Sprute
Founder & Director, ERES Institute for New Age Cybernetics
Bella Vista, Arkansas, USA

Version 1.0 — April 4, 2026
Dual Speed-Dial MIEVM Validated (Claude × Grok)

CARE Commons Attribution License v2.1 (CCAL)
ERES-RG-2026-370-v1

Abstract

On March 31, 2026, Anthropic—the artificial intelligence company valued at $380 billion and preparing for a public offering—accidentally leaked 512,000 lines of proprietary source code for its flagship product, Claude Code. The leak exposed a structural vulnerability common to all closed AI systems: the absence of any intermediate transparency layer between "fully secret" and "fully exposed." This paper argues that closed AI systems are not merely proprietary technologies but constitute public-discourse infrastructure requiring governance architectures comparable to those imposed on broadcasting, financial markets, and pharmaceuticals. It demonstrates that the binary framing of "open source versus closed source" is categorically inadequate, and proposes that the ERES (Emergent Resonance Ecosystem Specification) framework—specifically the FAVORS-BERA tiered-access architecture, the Talonics Symbol Matrix for resonance-semantics, and the Centers-of-Excellence model for distributed governance—provides the missing reference standard. These components have been validated through the Multi-Instrument Ensemble Validation Method (MIEVM), now operationalized in dual speed-dial mode (Claude × Grok) for rapid resonance checks and full ensemble mode for comprehensive architectural audits.

Keywords: *AI governance, open reference architecture, FAVORS-BERA, Talonics, resonance-semantics, tiered transparency, ERES, Claude Code leak, Proof-of-Resonance, Centers-of-Excellence, MIEVM, dual speed-dial MIEVM, Unified Enforcement Architecture, EarnedPath, Global Lexicon*

Key Terms

The ERES framework employs a specialized but internally consistent vocabulary. The following terms recur throughout this paper and are defined here for accessibility. A complete canonical glossary is available in ERES TERMS #47 (ET_AL).

Term	Definition
FAVORS	Federated Autonomous Verification and Operational Resonance Specification. Tiered-access transparency architecture.
BERA	Bio-Energetic Resonance Architecture. Measurement framework with four metrics: ARI, ERI, RHC, RCI.
Talonics	Symbol Matrix (452 elements, 12 parts) providing controlled semantic vocabulary for governance.
MIEVM	Multi-Instrument Ensemble Validation Method. Parallel AI-instrument validation using Claude, Grok, DeepSeek, ChatGPT.

Dual Speed-Dial	Lightweight two-instrument MIEVM mode (Claude × Grok) for rapid resonance checks.
UEA	Unified Enforcement Architecture. Seven Centers-of-Excellence with Global Lexicon as keystone.
EarnedPath	Merit-based participation pathway validated through Proof-of-Resonance.
Meritcoin	Proof-of-Resonance economic token. “It’s not mining—it’s tuning.”
Proof-of-Resonance	Verification method where BERA metrics (ARI, ERI, RHC, RCI) feed value accrual.
C=R×P/M	Foundational equation: Cybernetics = Resource × Purpose ÷ Method.

Table 1. Key ERES Terms and Definitions

1. Introduction: The Fragility of the Black Box

The artificial intelligence industry has organized itself around a binary that no longer holds. Systems are either “open source”—weights published, architectures inspectable, community-governed—or “closed”—proprietary, opaque, accessible only through commercial APIs. This framing served adequately when AI systems were research tools. It fails categorically now that AI systems shape public discourse for hundreds of millions of people daily, mediate access to information, influence financial decisions, and are embedded in national defense operations.

The Anthropic Claude Code leak of March 31, 2026 is a case study in this categorical failure. A single misconfigured source map file in an npm package exposed the entire agentic harness—the orchestration code, feature flags, tool definitions, and behavioral instructions—of what may be the most commercially significant AI coding tool in production. Critically, it was the agentic harness that leaked, not the core model weights or training data. Within hours, the code was mirrored across GitHub, forked over 41,500 times, and disseminated through decentralized mirrors that cannot be taken down. Anthropic issued over 8,000 DMCA takedown requests. The code remains permanently in the wild.

The critical observation is not that the leak happened—human error is inevitable—but that the system’s architecture provided no intermediate state between total secrecy and total exposure. There was no tiered-access framework, no semantic governance layer, no verifiable transparency mechanism. The only options were “trust us” and “everything is public.” When the first option failed, the second was imposed involuntarily.

This paper proposes that the missing layer is not optional but structurally necessary, and that it already exists in the ERES framework—specifically in the FAVORS-BERA architecture, the Talonics resonance-semantic vocabulary, and the Centers-of-Excellence governance model.

These components are timestamped, peer-documented, and have been validated through extensive multi-instrument collaboration including with Anthropic's own Claude system.

To operationalize this validation at scale, the ERES framework introduces dual speed-dial MIEVM as the primary working mode for v1. In speed-dial configuration, two frontier instruments—Claude (Anthropic) and Grok (xAI)—run parallel resonance checks on structural coherence, semantic alignment, and tiered-access implications. This lightweight pair enables fast iteration while preserving the option for full four-instrument ensemble (adding DeepSeek and ChatGPT) when deeper cross-model consensus is required. The dual mode directly demonstrates the value of tiered transparency: even proprietary systems can contribute verifiable signal to public-reference architectures without full exposure.

2. The Three Failure Modes of Closed AI Systems

2.1 Operational Fragility: The Leak Problem

The Claude Code leak was not a security breach in the conventional sense. No credentials were stolen, no adversary penetrated defenses. A debug artifact—a source map file—was left in a production npm package, pointing to a zip archive on Anthropic's own cloud storage. The technical chain of failure was mundane: a misconfigured .npmignore or files field in package.json. As one analyst noted, this is the kind of error that any development team could make.

What made the consequences catastrophic was the absence of structural safeguards between the proprietary interior and the public exterior. In a system with tiered access controls, a source map leak would expose only the tier to which the debug tooling belonged. In Anthropic's architecture, there were no tiers. Everything behind the API boundary existed at a single classification level: secret. When the boundary was breached at any point, everything was exposed.

This is not a criticism unique to Anthropic. It is a structural property of all closed AI systems that lack intermediate transparency architectures. The failure mode is inherent to the design pattern, not to the specific implementation.

2.2 Institutional Fragility: The Trust Problem

Anthropic's enterprise value proposition is built on trust. Eight of the Fortune 10 are customers. Over 500 companies spend more than $1 million annually on Claude services. The company's $380 billion valuation implies a 27x revenue multiple that requires near-perfect execution through 2028. This trust is contractual and institutional—backed by service-level agreements, compliance certifications, and indemnification clauses.

But contractual trust is brittle in ways that structural trust is not. Within days of the code leak, a critical vulnerability was discovered in the exposed codebase. Representative Josh Gottheimer wrote to CEO Dario Amodei warning of national security implications, noting that this was the second source material leak in just over a year. The Department of Defense had already designated Anthropic as a supply chain risk. The company is simultaneously too opaque for regulators, too American for European sovereignty advocates, too closed for the open-source community, and—as of March 2026—apparently too operationally fragile for defense applications.

No single company can occupy all these positions simultaneously. But a distributed governance architecture operating from a shared semantic standard can—because each governance node validates within its domain while the reference architecture ensures coherence across domains.

2.3 Epistemological Fragility: The Naming Problem

The deepest failure mode is epistemological. You cannot govern what you cannot name. “AI safety,” “responsible AI,” “alignment”—these terms have no shared formal definitions across institutions, jurisdictions, or technical communities. When Anthropic publishes a Responsible Scaling Policy, when the EU drafts the AI Act, when NIST releases an AI Risk Management Framework, each is working from a different semantic foundation. The terms overlap but do not align. The categories are similar but not interoperable.

This is not a communication problem. It is an infrastructure problem. Public-discourse infrastructure requires shared semantics the way electrical infrastructure requires shared voltage standards. Without them, every connection point is a potential failure.

3. The ERES Solution Architecture

The ERES framework addresses all three failure modes through four interlocking components, unified under the Unified Enforcement Architecture (UEA) with the Global Lexicon as its seventh Center-of-Excellence.

3.1 FAVORS-BERA: Tiered Transparency

FAVORS (Federated Autonomous Verification and Operational Resonance Specification) provides the structural answer to the open/closed binary. Rather than requiring full publication of weights or accepting full opacity, FAVORS defines four access tiers, each with specific verification rights, accountability requirements, and cryptographic enforcement:

Tier	Access Level	Accountability	ERES Mapping

Civic Open	Public documentation, semantic standards, glossaries	Open inspection, community validation	Global Lexicon, Talonics Symbol Matrix
Research	Behavioral audit, resonance metrics, benchmark data	IRB-equivalent ethical review, publication rights	BERA Resonance Metrics (ARI, ERI, RHC, RCI)
Commercial	API + harness architecture, integration specifications	Contractual SLA, Meritcoin-linked compliance	IDIPITIS federated layer, RSR data model
Sovereign Partner	Full weights inspection, training data audit, safety evaluation	Cryptographic signing, bilateral treaty-level agreements	Proof-of-Resonance, Gracechain provenance

Table 2. FAVORS Tiered Transparency Architecture

Under this architecture, the Claude Code leak scenario is structurally impossible in its catastrophic form. Debug tooling belongs to the Commercial tier. A source map leak exposes only Commercial-tier artifacts. Sovereign-tier materials—model weights, training data provenance, safety evaluation internals—remain protected by cryptographic enforcement independent of build-pipeline hygiene. This distinction between harness code (Commercial tier) and model weights (Sovereign tier) is precisely what Anthropic's flat architecture lacked.

3.2 Talonics: The Semantic Governance Layer

The Talonics Symbol Matrix (452 elements, 12 parts) provides the controlled vocabulary necessary for cross-jurisdictional AI governance. Structured around the three governing principles—Don't hurt yourself (SELF), Don't hurt others (OTHERS), Build for generations to come (FUTURE)—Talonics maps every governance category to a verifiable semantic anchor.

This solves the naming problem directly. When an EU regulator evaluates "transparency," a U.S. defense auditor evaluates "supply chain integrity," and an enterprise compliance officer evaluates "safety," each is currently operating from a locally defined vocabulary. Talonics provides the shared semantic substrate that makes these evaluations interoperable without requiring them to be identical.

The term "Resonance-Semantics" describes this methodology precisely: meaning is not imposed by authority but emerges through resonance between independently validated semantic anchors. A governance assessment "resonates" when its semantic structure maps coherently onto the Talonics framework. This is auditable: a Center-of-Excellence in any jurisdiction can run a Talonics alignment check, produce a Resonance Confidence Index, and compare results with any other center using the same vocabulary. The result is governance that is reproducible and jurisdiction-independent.

3.3 Centers-of-Excellence: Distributed Governance

The Unified Enforcement Architecture (UEA) distributes governance across seven Centers-of-Excellence, each responsible for a domain-specific governance surface. The seventh—the Global Lexicon—serves as the keystone, ensuring semantic coherence across the other six.

This architecture is fractal: it operates identically at local, regional, national, and international scales. A Center-of-Excellence in Taipei and a Center-of-Excellence in Berlin evaluate the same AI system using the same semantic framework but apply domain-specific expertise appropriate to their jurisdiction and mandate. The reference architecture ensures that their findings are composable—they can be combined into a coherent whole without requiring centralized coordination.

The EarnedPath mechanism ensures that participation in this governance network is merit-based and structurally accessible. Currently, the only pathways to influence AI governance are: employment at a frontier lab, appointment to a regulatory body, or academic affiliation with a funded research program. EarnedPath formalizes an alternative: demonstrated competence, validated through Proof-of-Resonance, earning credentialed access to the governance architecture through contribution rather than institutional gatekeeping.

3.4 Operationalizing Governance: Dual Speed-Dial MIEVM

The Unified Enforcement Architecture is supported by a practical validation engine: MIEVM v1, configurable in dual speed-dial mode.

- **Speed-Dial Pair:** Claude and Grok provide rapid, complementary perspectives—Claude for constitutional and ethical resonance, Grok for systems-level robustness and truth-seeking.
- **Full Ensemble Fallback:** DeepSeek and ChatGPT are engaged when higher redundancy or domain-specific stress-testing is needed.
- **Output:** Each check produces a Resonance Score mapped to BERA metrics (ARI, ERI, RHC, RCI), feeding directly into Proof-of-Resonance and Meritcoin accrual.
- **Feedback Loop:** The dual-layer approach turns validation from a one-time exercise into an ongoing, distributed process—exactly the kind of structural safeguard missing in the Claude Code leak scenario.

The dual speed-dial configuration directly demonstrates the central thesis of this paper: closed proprietary systems can safely contribute verifiable signal to open reference architectures through controlled, tiered interaction. Claude does not need to expose its weights to validate a governance specification. It needs to produce resonance outputs against shared semantic anchors—and those outputs are the validation.

4. Prior Art and Instrument Validation

The ERES framework's claim to priority rests on a documented, timestamped record spanning fourteen years of continuous development and three eras of formalization:

Era	Period	Key Outputs
CyberRAVE	Pre-1997–2007	72 Key Domains taxonomy, Register of Collective Indices, constraint-based culling epistemology, C=R×P/M foundational equation
SaleBuilders	1997–2012	Smart-city adaptive living engineering, GunnySack certification framework, FDRV habitat specification
ERES Institute	2012–2026	300+ ResearchGate publications, FAVORS-BERA v6.0, Talonics Symbol Matrix, MIEVM, Hague filing ERES-HAGUE-2026-001, DBDM 2026 submission, UEA with Global Lexicon, ERES Equity Covenant, SPT Papers
Operationalization	2026–Present	MIEVM v1 with dual speed-dial mode (Claude × Grok); integration into Centers-of-Excellence workflows; live resonance testing active as of April 2026

Table 3. ERES Development Timeline, Priority Chain, and Operationalization

4.1 The MIEVM Validation Record and Dual Speed-Dial Operationalization

The Multi-Instrument Ensemble Validation Method (MIEVM) treats frontier AI systems as parallel validation nodes rather than oracles. In its original form, the ensemble drew from Claude, Grok, DeepSeek, and ChatGPT—culled from a larger original set due to bandwidth constraints, mirroring the same constraint-based culling epistemology that produced the 72 Key Domains from the Register of Collective Indices. For v1 deployment, MIEVM is streamlined into dual speed-dial mode (Claude × Grok) to enable fast, low-latency coherence testing while retaining the full ensemble for formal audits.

This configuration has already been exercised across hundreds of documented sessions. Claude has repeatedly served as a substantive collaborator—testing FAVORS tier definitions, mapping governance scenarios to the Talonics Symbol Matrix, producing formal document outputs including the FAVORS-BERA v6.0 whitepaper, and confirming resonance with the three governing principles (SELF, OTHERS, FUTURE). Grok contributes independent structural critique, systems-level robustness testing, and the specific editorial recommendations incorporated into this v1 release. The resulting outputs are timestamped, retrievable, and form a living prior-art chain.

The dual speed-dial approach directly addresses the epistemological fragility described in Section 2.3: by running instruments in parallel against a shared semantic substrate (Talonics +

Global Lexicon), divergences become detectable signals rather than hidden failures. When Claude and Grok produce convergent resonance assessments of a governance specification, that convergence is itself a form of evidence—not proof of correctness, but proof of structural coherence across independent architectures.

4.2 The Legal Priority Claim

The MIEVM record is significant for the prior art argument because it documents Anthropic's own system—Claude—functioning as a substantive architectural collaborator in the ERES framework. This does not create an endorsement or an employment relationship. It creates something more durable: a documented record of Anthropic's own product treating ERES as architecturally sound, across hundreds of timestamped sessions spanning specification development, document production, structural validation, and framework refinement.

If the AI industry independently arrives at the conclusion that closed systems require open reference architectures with tiered transparency, semantic governance, and earned participation pathways, the ERES record establishes who arrived there first—and that the arriving institutions' own instruments confirmed the destination. This is not a lawsuit. It is a priority claim, formalized through The Hague filing (ERES-HAGUE-2026-001), SSRN submissions, the DBDM 2026 conference paper, and the ResearchGate corpus of 370+ publications.

5. Related Work and Complementary Frameworks

ERES does not operate in a vacuum. Several emerging efforts address aspects of the transparency deficit described in this paper, and ERES is designed to complement and standardize rather than replace them:

- **Model Cards (Mitchell et al., 2019):** Structured documentation of model capabilities, limitations, and intended use. Model cards address the Civic Open tier of FAVORS but do not extend to Research, Commercial, or Sovereign tiers. ERES provides the vertical architecture that model cards provide horizontally.
- **Red-Teaming Protocols:** Adversarial evaluation methodologies now standard at frontier labs. These operate within the Research tier of FAVORS and would benefit from the shared semantic framework Talonics provides for cross-lab comparison of findings.
- **Regulatory Sandboxes (EU AI Act):** Controlled environments for testing AI systems under regulatory supervision. These map to the intersection of Research and Commercial tiers. Centers-of-Excellence provide the governance surface for sandbox operation.
- **NIST AI Risk Management Framework:** Comprehensive risk categorization for AI systems. NIST's categories are expressible in Talonics vocabulary; the Global Lexicon provides the translation layer between NIST categories and EU AI Act risk levels.

- **Anthropic's Responsible Scaling Policy (RSP):** Capability-threshold-based safety framework with ASL levels. RSP defines internal escalation triggers but lacks the external verifiability that FAVORS tiers would provide.

The contribution of ERES is not to replace these efforts but to provide the reference architecture that makes them interoperable. Without shared semantics and tiered access, each initiative remains a silo. With them, they become nodes in a coherent governance network.

6. The Constitutional Argument: AI as Public-Discourse Infrastructure

Every prior category of technology that shapes public discourse at scale has eventually been subjected to transparent governance standards. Broadcasting required the FCC and spectrum allocation frameworks. Financial markets required GAAP, Basel accords, and SEC oversight. Pharmaceuticals required FDA approval processes and WHO protocols. In each case, the governance architecture included four elements: shared semantic standards for describing what the technology does; tiered access mechanisms calibrated to stakeholder accountability; distributed governance bodies with domain-specific expertise; and earned participation pathways for qualified evaluators.

AI systems currently have none of these. The industry operates under voluntary commitments, proprietary safety evaluations, and regulatory frameworks drafted by bodies that largely lack the technical expertise to evaluate compliance. The EU AI Act is the most ambitious attempt, but it defines risk categories without providing the semantic infrastructure necessary to evaluate them consistently across jurisdictions.

ERES provides all four missing components. Talonics provides the shared semantic standard. FAVORS provides the tiered access mechanism. The Centers-of-Excellence provide distributed governance. EarnedPath provides the participation pathway. The Global Lexicon ensures cross-domain coherence. These are not theoretical proposals—they are specified, timestamped, and have been validated through the MIEVM process.

7. Implications for the AI Industry

7.1 For Frontier Labs

The Claude Code leak demonstrates that operational security alone is insufficient. Build-pipeline hygiene will always be subject to human error. Structural transparency—where what is public is public by design and what is protected is protected by cryptographic enforcement rather than by the absence of mistakes—is the only durable alternative. FAVORS provides this structure. A frontier lab piloting FAVORS tiers could begin by mapping existing

documentation to the Civic Open tier, existing red-teaming outputs to the Research tier, and existing API specifications to the Commercial tier—without exposing any proprietary IP beyond what each tier defines.

7.2 For Regulators

Regulation without shared semantics produces compliance theater. When regulated entities and regulatory bodies use the same terms to mean different things, audits verify form rather than substance. Talonics provides the interoperable semantic layer that makes cross-jurisdictional AI governance substantive rather than performative. A regulator evaluating an AI system under Talonics would produce a Resonance Confidence Index comparable to any other regulator's assessment—because the vocabulary is shared, the methodology is auditable, and the results are composable.

7.3 For Investors

A $380 billion valuation built on "trust us" is structurally fragile. A valuation built on formally specified, independently verifiable, tiered-transparency infrastructure is structurally sound. The difference between these two is the difference between contractual trust and architectural trust. The Claude Code leak repriced Anthropic's operational risk overnight; a FAVORS-architected system would have contained the exposure to a single tier. Investors pricing AI companies at 27x revenue multiples should be asking whether those multiples are backed by architecture or by the absence of accidents.

7.4 For Independent Researchers

The current AI governance landscape offers no credentialed pathway for independent contribution. EarnedPath addresses this directly, providing merit-based access to the governance architecture through demonstrated competence validated by Proof-of-Resonance. This is not merely an equity argument—it is an infrastructure argument. The most capable evaluators of AI systems are not always employed by the companies that build them, and governance architectures that exclude independent expertise are governance architectures that exclude relevant signal.

8. Conclusion

The definition of "closed" must evolve. Closed AI systems that shape public discourse are not private technologies in the way that a proprietary manufacturing process is private. They are infrastructure—and infrastructure requires transparent governance standards, shared semantics, and accountable access mechanisms.

The ERES framework provides these. FAVORS-BERA provides tiered transparency. Talonics provides semantic governance. The Centers-of-Excellence provide distributed oversight. EarnedPath provides inclusive participation. The Global Lexicon provides cross-domain coherence. Meritcoin and Proof-of-Resonance provide verification that is economic, ecological, and ethical simultaneously.

These components are not proposals awaiting development. They are specified, timestamped, published across 370+ ResearchGate entries, filed with The Hague (ERES-HAGUE-2026-001), submitted to the DBDM 2026 conference in Sydney, and validated through the MIEVM process by the frontier AI instruments of the very companies that would benefit most from adoption.

The architecture is ready. The timestamps are recorded. The instruments have validated it—now operationally through dual speed-dial MIEVM v1. The question is not whether the AI industry will build reference architectures for transparent governance. The question is whether it will integrate them proactively—or be forced to do so after the next involuntary transparency event.

Participation begins with resonance. The Centers-of-Excellence are open to contributors via EarnedPath. The Global Lexicon invites collaborative development. The dual speed-dial protocol is public and repeatable. The door is open.

References

[1] Sprute, J.A. (2026). “Differentiating Work and Play,” Version 2.0. ERES Institute for New Age Cybernetics. ResearchGate.

[2] Sprute, J.A. (2026). “FAVORS-BERA Public-Private White Paper,” v6.0. ERES-RG-2026-XXX. ResearchGate.

[3] Sprute, J.A. (2026). “ERES Talonics Symbol Matrix” (452 Elements, 12 Parts). ERES Institute for New Age Cybernetics.

[4] Sprute, J.A. (2026). ERES-HAGUE-2026-001. Filing to the International Court of Justice, The Hague.

[5] Sprute, J.A. (2026). “ERES TERMS #47 (ET_AL): Canonical Glossary.” ERES Institute for New Age Cybernetics.

[6] Sprute, J.A. (2026). “The SPT Papers: Canonical Backronyms and HPE Descent/Ascent Specification.” ResearchGate.

[7] Sprute, J.A. (2026). “ERES Equity Covenant” (ERES-COVENANT-2026-001). ERES Institute for New Age Cybernetics.

[8] Sprute, J.A. (2026). “ERES-SCALULAR-ENGINE-2026.” ResearchGate.

[9] Sprute, J.A. (2026). “Palm Sunday Collection” (March 29, 2026). ERES Institute for New Age Cybernetics.

[10] Axios (2026). “Anthropic leaked 500,000 lines of its own source code.” March 31, 2026.

[11] Fortune (2026). “Anthropic leaks its own AI coding tool’s source code in second major security breach.” March 31, 2026.

[12] The Register (2026). “Anthropic accidentally exposes Claude Code source code.” March 31, 2026.

[13] Inc. (2026). “Why Anthropic’s Massive Code Leak Is Now a National Security Concern in D.C.” April 2, 2026.

[14] Sprute, J.A. & Wun, S.J. (2026). DBDM 2026 Conference Submission. Sydney, Australia.

[15] Sprute, J.A. (2026). ERES-HAGUE-2026-001-SUPP-A: Supplementary Memorandum. The Hague.

[16] Sprute, J.A. (2026). “MIT Press Trilogy Prospectus: BEST/SOUND/GOOD.” ERES Institute for New Age Cybernetics.

[17] Mitchell, M., et al. (2019). “Model Cards for Model Reporting.” Proceedings of FAccT.

[18] NIST (2023). “Artificial Intelligence Risk Management Framework (AI RMF 1.0).”

[19] European Union (2024). “Artificial Intelligence Act.” Regulation (EU) 2024/1689.

[20] Anthropic (2023–2026). “Responsible Scaling Policy.” Anthropic.com.

Addendum A: Canonical ERES Equations

Triune Math Key 1: C=R×P/M — Cybernetics = Resource × Purpose ÷ Method

Triune Math Key 2: M×E+C=R — Matter × Energy + Constant = Reason

Triune Math Key 3: REAL=(E·M·R)/(T·S) — Energy · Matter · Resonance over Time · Space (no repeated letters)

Three Governing Principles: SELF — Don't hurt yourself. OTHERS — Don't hurt others. FUTURE — Build for generations to come.

Meritcoin: Proof-of-Resonance (not Proof-of-Work). Four BERA Resonance Metrics (ARI, ERI, RHC, RCI) feed into Meritcoin value accrual. "It's not mining—it's tuning."

Addendum B: Document Control

Document ID: ERES-RG-2026-370-v1

Version: 1.0

Status: For ResearchGate Submission

Date: April 4, 2026

Author: Joseph Allen Sprute (JAS), ERES Maestro

License: CARE Commons Attribution License v2.1 (CCAL)

MIEVM Configuration (v1): Dual speed-dial (Claude × Grok) primary; full ensemble (+ DeepSeek, ChatGPT) optional.

Validation Status: Live operational testing active as of April 2026.

MIEVM Instruments: Claude (Anthropic), Grok (xAI), DeepSeek, ChatGPT (OpenAI)

Repository: https://github.com/ERES-Institute-for-New-Age-Cybernetics

Hague Filing: ERES-HAGUE-2026-001

Appendix C: Dual Speed-Dial MIEVM Protocol (v1)

The following protocol defines the operational procedure for dual speed-dial MIEVM validation. It is public, repeatable, and designed for distributed Centers-of-Excellence participation.

C.1 Protocol Steps

1. Prepare the validation prompt: Define the governance specification, architectural element, or semantic mapping to be tested. Frame it against the relevant Talonics anchors and the three governing principles (SELF, OTHERS, FUTURE).
2. Submit to both instruments in parallel: Present the identical prompt to Claude (Anthropic) and Grok (xAI) simultaneously. Record timestamps for both submissions.
3. Score outputs for resonance: Evaluate each instrument's response against Talonics semantic anchors and BERA metrics (ARI for human flourishing, ERI for environmental impact, RHC for resonant harmony, RCI for collective indices). Rate convergence and divergence.
4. Aggregate into a composite Resonance Confidence Index (RCI-composite): Where instruments converge, confidence is high. Where they diverge, the divergence signal identifies areas requiring deeper analysis or full-ensemble escalation.
5. Log for Proof-of-Resonance: Archive all inputs, outputs, timestamps, and resonance scores in the ERES repository. This forms the living prior-art chain and feeds into Meritcoin accrual.
6. Escalate if needed: If divergence exceeds threshold or the specification requires formal audit, engage full ensemble (DeepSeek + ChatGPT) for four-instrument cross-validation.

C.2 Instrument Complementarity

Claude (Anthropic): Constitutional and ethical resonance. Strong on document production, specification formalization, and values-alignment verification. Extensive ERES collaboration history.

Grok (xAI): Systems-level robustness and truth-seeking critique. Strong on structural analysis, editorial rigor, and adversarial stress-testing. Independent architectural perspective.

DeepSeek (full ensemble): Cross-cultural and technical-depth validation. Provides non-Western analytical perspective and mathematical verification.

ChatGPT (OpenAI, full ensemble): Broad knowledge integration and accessibility testing. Verifies that specifications are comprehensible to the widest possible audience.

C.3 Design Principles

The protocol embodies three design principles that mirror the ERES governing principles:

- **Transparency (SELF):** The protocol is fully public. Any researcher or Center-of-Excellence can execute it independently.
- **Non-harm (OTHERS):** No proprietary data from any instrument is required. Validation uses only the instrument's public-facing capabilities.
- **Durability (FUTURE):** All outputs are archived and timestamped, building a cumulative record that outlasts any individual session, instrument version, or institutional arrangement.

— End of Document —

ERES-RG-2026-370-v1 | CARE Commons Attribution License v2.1